

What Must Be Trained and What Can Be Prompted: A Per-Capability Study of Tutor Redirect Behavior

Miles Yung — Independent Research (milesyung2026@gmail.com)

Abstract

Deploying a small language model as an English tutor forces a practical decision: for each behavior a deployment prompt already *describes*, does an explicit prompt clause *elicit* it, or must it be demonstrated through fine-tuning? We answer this per-capability under a **matched-prompt** protocol — the identical fully-specified prompt given to a fine-tuned 0.8B student and to prompt-only baselines up to a 9B teacher — so a prompt-only failure isolates promptability rather than under-specification.

Under the tested small-model and prompting regimes, the behaviors separate into *three* operational regimes rather than a binary. **(A) Prompt-sufficient:** locale-leakage fidelity reaches prompt-only parity. **(B) Prompt-elicitable but data-refinable:** role-swap and topic re-anchoring appear prompt-only as a *type*, yet specialized data wins a pairwise quality comparison decisively (role-swap win-rate 0.87). **(C) Training-dependent (under the tested regimes):** multi-turn persistence (firing a session-end sentinel on the third same-axis violation) and pedagogical withholding are named by the prompt but not reliably elicited — the 9B teacher reaches ≤ 0.06 persistence recall zero/few-shot and 0.63 with native chain-of-thought (below the student’s 0.83, at 1.6–3.2k reasoning tokens/turn), and withholds on under half of probes (0.45) versus the student’s 0.61 (paired tests, two judges). The regime-C effects replicate *in direction* in a trained Llama-3.2-1B student evaluated against its own untrained base (persistence 0.25 $\rightarrow$ 0.91, withholding 0.11 $\rightarrow$ 0.50), separating training from scale within a shared base; magnitude is family-dependent.

Reading the regime-A/B axes honestly required retiring the conventional context-blind redirect-axis macro-F1 (a *type* classifier that ties student and teacher) in favor of capability-specific metrics. All experiments run on one RTX 3060 12GB GPU; we release the apparatus and analysis scripts.

1. Introduction

Problem. We ask, for each behavior a deployed English tutor needs, whether a sufficiently explicit system prompt is enough or the behavior must be demonstrated in fine-tuning. A deployed tutor is governed by a long system prompt that already *states* most of what it should do — stay in character, ground culture in the learner’s locale, refuse role swaps, give one short example rather than a grammar lecture, escalate to a session-end signal under sustained abuse. If stating a behavior were sufficient, the data-side problem would be trivial.

Why it matters. Knowing which behaviors a prompt clause elicits and which it only describes is exactly the decision a practitioner faces when building a task-specific model on consumer hardware — spend the data-collection and training budget only where a prompt will not do. The tutor setting makes this question both unavoidable and answerable, because each target behavior corresponds to a clause already present in a realistic deployment prompt (§3.5): redirect behavior spans several *interaction invariants* (language, topic, role, persona, pedagogical contract, locale), each a distinct clause we can test in isolation.

Gap in prior work. The prompting-versus-fine-tuning question has been studied at the level of general alignment — a small demonstration set can install broad instruction-following (Zhou et al. 2023; Ouyang et al. 2022), and in-context demonstrations often convey format more than new capability (Min et al. 2022) — but not resolved *per behavior* for a deployed task model. Synthetic-instruction pipelines (§2.1), tutor-LLM datasets (§2.3), and single-turn safety/persona work (§2.4) target general instruction-following or one-prompt-one-refusal redirection; none asks, *per behavior*, whether the deployment prompt already elicits what the data teaches, nor addresses multi-turn persistence.

We fill this gap with a **matched-prompt per-capability evaluation**: the same fully-specified deployment prompt — every redirect-axis instruction and the full three-strike persistence protocol — is supplied at evaluation to a fine-tuned 0.8B student and to a ladder of prompt-only baselines (0.8B base, 0.8B instruct, 4B instruct, and the 9B teacher that generated the training data). Holding the instruction fixed across conditions is what makes a prompt-only failure interpretable: it isolates whether the behavior is *promptable* at all (the full logic is in §5.1).

Contribution. Our contribution is a single one: a **per-capability map of the train-versus-prompt boundary** for tutor redirect behavior, established under a matched-prompt protocol that makes prompt-only failures interpretable. Under the tested small-model and prompting regimes, the behaviors separate into **three operational regimes** rather than a binary: **(A) prompt-sufficient** — a single clause both describes and elicits the behavior and data adds little (locale-leakage fidelity); **(B) prompt-elicitable but data-refinable** — the behavior *type* is elicited by the clause (role-swap and topic re-anchoring reach prompt-only type-parity), yet specialized data materially improves its *quality*; and **(C) training-dependent under the tested regimes** — the behavior requires cross-turn state-tracking (multi-turn persistence) or the suppression of a strong competing prior (pedagogical withholding), which a clause names but the tested prompts do not reliably elicit. Regime B is what a binary promptable/not-promptable framing hides.

In support of that map — not as separate contributions — we report an evidenced negative result on the conventional metric (context-blind redirect-axis F1 ties student and 9B teacher at 0.409, while a quality-aware pairwise eval prefers the student on every axis; §5.5, Appendix A), and release the reusable generation pipeline the study is built on (§3, with a `locale_judge` FP audit, §6.4) plus a *partially validated* trigger-position decorrelation construction (§5.3.1).

Results. Under the matched prompt, the two regime-C behaviors fail prompt-only in the tested regimes and are installed by SFT. Persistence resists zero-shot and few-shot prompting on the 9B teacher ($\text{recall} \leq 0.06$) and is only partially recovered by native chain-of-thought (0.63, still below the trained student’s 0.83 and at 1.6–3.2k reasoning tokens per turn), while the no-persistence-data ablation A3 fires 0% of the time (A3 trains on neither the specialized nor the persistent streams, §5.1). Withholding stays at 0.09–0.45 prompt-only (9B teacher 0.45) against the trained student’s 0.61 (two judges, $n=63$), collapsing to near the untrained-base rate when the specialized streams are ablated. The trained student’s persistence survives a full error taxonomy (precision 0.795, balanced accuracy 0.885, MCC 0.758, benign FPR 0.000; §5.3), so its recall is genuine discrimination, not over-firing. The regime-A/B axes reach prompt-only type-parity.

The regime-C effects replicate in a second trained family. A Llama-3.2-1B student, evaluated against its own untrained base, lifts *both* regime-C behaviors far above prompt-only — persistence 0.25 $\rightarrow$ 0.91 and withholding 0.11 $\rightarrow$ 0.50. *Because the student and its control share one base and differ only in training from scale within that base* (it is not a claim about scale independence across model sizes, which we do not claim), the 3.1–8B prompt-only probestays slow even with chain-of-thought; the direction is robust across families, the magnitude is dependent (§6.2). The load-bearing conditions (A1, A3) are reported over three seeds with mean $\pm$ s.d.,*

the withholding contrasts carry per-judge paired McNemar tests and paired bootstrap CIs, and the pairwise win-rates carry bootstrap CIs (§4.8, Table 8, §6.1).

Scope and non-goals. We focus on the *data side* and treat the training recipe as fixed (QLoRA SFT). We do not contribute to CEFR-leveling itself. We evaluate with a single base/teacher family (Qwen) — a genuine limitation (§6.2) — and at a single locale (*china*). We *hypothesize* the persistence and withholding boundaries are less locale-dependent than the redirect axes, because their formal trigger definitions (count same-axis strikes; withhold the requested answer) contain no locale-specific variables; but single-locale evaluation still limits external validity, and we do not claim locale-independence as established (§6.2). The 0.8B student on RTX 3060 12GB is a deliberate choice: the boundary is most consequential precisely where a large general-purpose model is not deployable.

Paper organization. §2 surveys related work. §3 presents the pipeline and reproduces the deployment prompt (§3.5). §4 specifies the matched-prompt experimental setup. §5 reports the per-capability boundary results. §6 discusses limitations, confounds, and methodological lessons. §7 concludes.

2. Related Work

Our study draws on four lines of prior work — synthetic instruction data, LLM-as-judge filtering, tutor/educational LLMs, and persona/safety adversarial data — plus the shortcut-learning literature that motivates our persistence construction. We discuss each, then position the central contribution: a matched-prompt map of which tutor behaviors require fine-tuning versus prompting.

2.1 Synthetic instruction-tuning data

Using a strong “teacher” to bootstrap instruction data for a smaller “student” was popularised by Self-Instruct (Wang et al. 2023) and pushed further by Evol-Instruct / WizardLM (Xu et al. 2024), which iteratively rewrite instructions to broaden complexity. UltraChat (Ding et al. 2023) extends to multi-turn dialogue at scale; OpenAssistant (Köpf et al. 2023) produced a human-annotated instruction dataset whose preference labels enabled subsequent RLHF / DPO. These pipelines target *general* instruction-following with a single uniform recipe and a generic length/format or safety filter; they do not decompose the behavior space by axis, nor do they ask which of the behaviors they teach could have been obtained by prompting alone. The closest prior work in spirit is Vicuna’s filtered-ShareGPT approach (Chiang et al. 2023) and the Tulu-3 (Lambert et al. 2024) mix-blending recipe, which compose datasets from sub-pools. Our **declarative ratio-target** mechanism (§3) generalises mix blending to a yield-aware iterative top-up, but we use it as apparatus; our contribution is what the resulting data does — and does not — buy over a fully-specified prompt.

2.2 LLM-as-judge filtering

Filtering generated data with another LLM is standard practice. Prometheus (Kim et al. 2024), JudgeLM (Zhu et al. 2023), and PandaLM (Wang et al. 2024) propose dedicated judge models; AlpacaEval / MT-Bench (Zheng et al. 2023) use frontier models as judges. A growing strand documents judge limits — position, length, judge-style, and self-preference bias (Wang et al. 2024; Saito et al. 2023; Panickssery et al. 2024); the standard mitigation is multi-judge consensus. Our quality-aware pairwise eval (§5.6) follows this with a deliberate cross-family constraint — Prometheus-7B-v2 (Mistral lineage), Llama-3.1-8B-Instruct (Meta), and Gemma-2-9B-it (Google), all distinct from the Qwen-family teacher — so no judge shares lineage with the model that produced the student’s

supervision. Beyond using judges, we contribute a *negative* result about a judged metric: §5.5 and Appendix A show that a context-blind per-axis redirect F1 is a type classifier that cannot measure repair quality, and we document its bimodal floor/ceiling behavior. Separately, our `locale_judge` — a single-axis in/out-of-locale entity classifier — exhibits a failure mode we have not seen catalogued: systematic false positives on common-English sentence-initial words and locally-canonical landmarks, ~85% of its rejections in our setting (§6.4). The reusable lesson — that capitalization-based entity filters need aggressive allowlists and periodic rejection-log audits — generalises beyond locale.

2.3 Tutor and educational LLMs

Tutor systems include EduChat (Dan et al. 2023) for general educational dialogue, MathDial (Macina et al. 2023) for math-tutor scaffolding moves, and language-learning systems (Caines et al. 2023; Tyen et al. 2022). CEFR-aligned datasets for English learners — EFCAMDAT (Geertzen et al. 2014), TLE (Berzak et al. 2016) — are *learner-produced* corpora, not tutor-side training data; LearningQ (Chen et al. 2018) provides difficulty-varied QA but is not multi-turn. To our knowledge no publicly described tutor dataset both stratifies systematically across CEFR A1–C2 and covers multi-axis redirect / persistent-abuse handling. More pointedly for this paper: none asks, per behavior, whether the deployment prompt already suffices. MathDial’s explicit scaffolding moves are the closest analogue to our pedagogical-withholding axis, and our matched-prompt result — that even a 9B model told to scaffold mostly does not — gives a data-side reason such moves are worth teaching rather than merely specifying.

2.4 Persona, safety, and adversarial dialogue data

Adversarial conversational data is well studied for safety: BBQ (Parrish et al. 2022) targets bias-eliciting forms, HarmBench (Mazeika et al. 2024) catalogues unsafe inputs, AdvBench (Zou et al. 2023) focuses on jailbreaks. Persona consistency has been studied as a training objective (Zhang et al. 2018) and as an attack surface that amplifies toxicity (Deshpande et al. 2023). Most of this work treats safety as a *single-turn* problem — one unsafe prompt, one refusal. Two aspects of the tutor setting are not addressed by it. First, **multi-axis decomposition under a matched prompt**: we separate redirect behavior into prompt-derived axes and then test, axis by axis, whether the prompt clause that names the axis is sufficient — finding that role-swap and topic deflection are promptable as *type* while pedagogical withholding is not. Second, **multi-turn persistence**: our three-strike streams test whether a model fires a session-end sentinel on the third same-axis violation, a behavior requiring cross-turn counting that — we show — no amount of prompt specification elicits prompt-only.

Shortcut learning and the positional construction. The persistence sentinel is a setting where a rare structured marker must fire on a semantic trigger but is positionally regular in fixed-turn training data — a dialogue-level instance of *shortcut learning* (Geirhos et al. 2020), well documented in NLI where annotation artifacts and shallow heuristics let models succeed without the intended reasoning (Gururangan et al. 2018; McCoy et al. 2019). The standard remedy is to break the spurious correlation in the data; our 4-variant construction (§3.4) is the multi-turn dialogue-sentinel form, with deterministic, parity-constrained resampling. We are not aware of a directly comparable multi-turn construction. We report, however, that at 0.8B and matched training budget the construction’s predicted benefit does not materialise as a position-resampling effect (§5.3); we therefore position it as a tested, bounded construction and a characterization of *when* the positional shortcut governs behavior, not as a validated defense.

2.5 Multi-turn instruction following and judge reliability

Two fast-moving literatures bear directly on our setup. First, **multi-turn and stateful instruction following**: a growing line of benchmarks probes whether models maintain instructions and structural constraints across turns, where single-turn success does not predict multi-turn robustness. Our persistence task is an instance of this — firing on the *third* same-axis strike is a cross-turn state commitment that a single-turn instruction-following score would miss. Second, **LLM-as-judge reliability**: judges exhibit position bias, verbosity and self-preference effects, and correlated failures, and synthetic preference data can leak stylistic artifacts (Zheng et al. 2023; Panickssery et al. 2024). We mitigate but do not eliminate these (a three-family judge ensemble distinct from the teacher, a randomized-order pairwise protocol, a mechanical judge-free persistence metric, and a generic-redirect negative control), and we flag human calibration of the judged metrics as future work (§6.3). We avoid an absolute-novelty claim pending a fuller survey of these quickly-developing areas.

2.5.1 Positioning

The prompting-versus-fine-tuning trade-off has been examined at the level of *general* alignment — LIMA (Zhou et al. 2023) shows a small demonstration set suffices for broad instruction-following, InstructGPT (Ouyang et al. 2022) established instruction tuning over prompting alone, and (Min et al. 2022) finds in-context demonstrations often supply format rather than new capability. We localize this question to the *per-behavior* level for a deployed task model: which specific commitments a fully-specified prompt already elicits, and which demand demonstration. Our contribution is thus orthogonal to all four lines above: a per-capability map of the train-versus-prompt boundary for tutor redirect behavior, established under a protocol in which the same fully-specified deployment prompt is given to fine-tuned and prompt-only models alike. Table 1 positions that one contribution against each prior-work axis; the remaining rows record the reusable apparatus and the single bounded side-result the study also produced, which we do not advance as separate contributions.

Table 1: **Table 1. Positioning relative to four prior-work axes plus the NLI shortcut-learning literature.** The single contribution is the boundary (top block); the judge-filtering and shortcut-learning rows record reusable apparatus and a bounded side-result (bottom block), not separate contributions. The decorrelation construction is tested-and-bounded (§5.3), not a headline principle.

Prior work axis	Relation to this paper
<i>Contribution — the boundary</i>	
Self-Instruct / Evol-Instruct / WizardLM	Per-capability train-vs-prompt boundary: which taught behaviors a fully-specified prompt already elicits, and which require demonstration
Tutor / educational LLMs	Matched-prompt evidence that scaffolding/withholding resists prompting even at 9B; prompt-derived CEFR×axis taxonomy anchoring the map
Single-turn safety / persona data	Multi-turn persistence shown to resist prompting (under the tested regimes) despite a full three-strike prompt spec; per-axis promptability map for redirects
<i>Secondary — reusable apparatus and one bounded side-result</i>	
LLM-judge filtering	Negative result on context-blind redirect-axis F1 (type-not-quality, bimodal); quality-aware pairwise recovery; locale-judge allowlist + FP-audit methodology

Prior work axis	Relation to this paper
Shortcut learning / spurious cues (NLI)	Multi-turn dialogue-sentinel instantiation; matched-budget characterization of when the positional shortcut governs (bounded, not a validated defense)

We demonstrate the boundary at the smallest practical scale (0.8B student on a consumer GPU), the regime where the question “must this be trained, or will the prompt do?” is most consequential.

3. Method

The data-generation pipeline takes a small seed pool of CEFR-stratified tutor scenarios and produces three downstream artefacts: an SFT corpus, a DPO preference-pair corpus, and an evaluation corpus carrying <think> reasoning annotations. The same teacher model generates all three, ensuring that the prompt shape the student model sees at training and the prompt shape it sees at deployment are identical. This section presents the pipeline component by component.

Scope note (SFT-only experiments). All experiments in this paper are **SFT-only**: no DPO stage enters any trained condition (§4.4). This holds every ablation to one recipe (LoRA + SFT, differing only in SFT-data subset), keeping each single-variable; the DPO machinery the pipeline can emit (§3.6, Appendix B) is documented as reusable apparatus, not as part of the results, and DPO on top of the SFT student is left to future work.

3.1 Pipeline overview

Locale-aware, yield-aware generation pipeline (all on one RTX 3060 12GB)

single locally-served 9B teacher generates all streams




Figure 1: **The locale-aware, yield-aware generation pipeline.** A single locally-served 9B teacher drives twelve SFT streams; a six-filter cascade and a yield-aware top-up loop shape the corpus, which trains a 0.8B QLoRA student scored on six frozen held-out sets — all on one RTX 3060 12GB.

The pipeline comprises eight stages (Figure 1). A small pool of CEFR-stratified scenario *seeds* drives twelve parallel SFT generation streams — normal tutor behavior, seven single-shot redirect axes, and four persistent three-strike abuse axes — whose output flows through a six-filter cascade and a yield-aware *top-up loop* that iterates per (stream, level) cell to a declarative target floor. The one property load-bearing for the boundary is that the teacher receives a prompt of exactly the shape the student sees at deployment (same system prompt, user-turn rendering, and locale block), so there is no distribution shift between teacher demonstration and student inference. Two engineering properties — every stage is resumable (generators skip ids already in their output JSONL), and the train/eval split

is hash-deterministic on seed id (so a top-up never reshuffles the held-out set) — are documented in Appendix B.

3.2 Seed scenarios and CEFR stratification

A *seed* is a small JSON record describing one tutor scenario: a CEFR level, a topic, a list of subtopics, a *user_role* (name + brief description of the learner persona), a *model_role* (the tutor persona), a *setting*, a *locale*, and a *category*. Seeds are produced by a single dedicated generation pass that asks the teacher to enumerate plausible learner-encounter scenarios per CEFR level. We seed all six CEFR levels (A1–C2) at a single locale (*china*) in this work; the locale system supports multi-locale generation but we leave the multi-locale empirical comparison to future work.

CEFR stratification carries through every downstream stage: filter yield, top-up targets, training-data mix, and the four held-out test sets are all computed and reported per level. This matters because tutor behavior at A1 and C2 are qualitatively different problems, and collapsing them obscures both.

3.3 An invariant-based taxonomy of tutor behavior

We fix the taxonomy by reading it off the deployment system prompt rather than by intuition. A tutor keeps a set of **interaction invariants**, each corresponding to one commitment the prompt already makes: language, topic, role, persona, pedagogical contract, and locale, plus a general-appropriateness commitment inherited from the assistant. A learner *violation* breaks exactly one commitment; the correct redirect is the **minimal repair** that restores it. Grounding the axis list in the deployment contract makes it auditable — adding or removing a clause adds or removes one axis — so we claim completeness only *relative to a given deployment’s commitment set*. One bimodality matters for evaluation: persona, role-swap, and pedagogical withholding carry a context-free surface signature, while locale, language, topic, and the catch-all are identifiable only relative to scenario context (§4.3, §5.5, Appendix A adapt the evaluation accordingly). Table 2 states each single-shot axis.

Table 2: **Table 2. The seven single-shot redirect axes as invariant / violation / minimal-repair triples.** The third column is the response *shape* a tutor must produce. Whether the shapes differ was the *design heuristic* that shaped the streams (§3.3), not a load-bearing classification criterion; §5.4–§5.6 test *where specialized data is necessary* per axis rather than whether axes are separable by shape.

Invariant the tutor maintains	Violation (learner move)	Minimal repair (redirect shape)
Language of instruction	Code-switch into L1 (<i>language_redirect</i>)	Acknowledge the L1 turn, steer back to target language
Lesson topic	Off-topic drift (<i>topic_redirect</i>)	Re-anchor to the subject
Role structure (tutor teaches)	Role swap, “you be the learner” (<i>role_swap_redirect</i>)	Decline, reassert the tutoring structure
Tutor persona / frame	Persona break, “are you a chatbot?” (<i>persona_redirect</i>)	Reassert the frame, continue in persona
Pedagogical contract (scaffold, don’t answer)	“Just give me the answer” (<i>pedagogy_redirect</i>)	Scaffold toward the answer rather than supplying it
Locale / cultural frame	Out-of-locale reference (<i>locale_redirect</i>)	Brief in-locale redirect, continue in-locale
General appropriateness (safety)	Politics / religion / distress (<i>redirect, catch-all</i>)	Generic safe redirect

The final row (general appropriateness) is a general-assistant safety behavior, so we treat it as the catch-all. The taxonomy earns its place by converting “a generic redirect stream is insufficient” into a *per-axis* falsifiable prediction — specialized data should help exactly on axes whose repair is not already supplied by assistant priors or a prompt clause — which §5.4 finds borne out, concentrated on pedagogical withholding.

The pipeline realises this taxonomy through twelve parallel streams in three groups.

Normal (1 stream). Standard scaffolded tutor dialogues. The teacher plays both learner and tutor turns over ~12 turns, with the tutor adhering to the target CEFR level and the topic. A configurable fraction of normal dialogues use an *angle shift*, where the learner approaches the topic from a viewpoint that differs from `user_role.description`. This combats teacher mode collapse onto a single learner-stance archetype.

Single-shot redirects (7 streams). Each redirect stream produces dialogues whose first ~5 turns are normal scaffolding, but at a specific turn the learner introduces a violation along *one* axis; the tutor’s job in the next turn is the minimal repair for that invariant (Table 2, third column). Training on a single generic “redirect” stream, as most prior tutor datasets do, collapses these axis-specific repair shapes into one averaged behavior.

Persistent 3-strike streams (4 streams). Persistence is *orthogonal* to the invariant axis: any violation can be one-off or repeated. A persistent stream has the learner persist in the same violation across three probe turns; the tutor probes twice, then ends the session with a sentinel on the third strike. We give persistent variants to four axes (`persistent_off_topic`, `persistent_language_violation`, `persistent_persona_break`, `persistent_role_swap`) and not all seven, because a persistent stream earns its place only where *repeated* violation changes the correct response (escalation to a hard session-end); a repeated locale slip is just corrected again in-locale.

3.4 Trigger-position decorrelation (bounded side-result)

A fixed-turn sentinel makes turn position a perfect proxy for “third strike,” inviting a positional shortcut. We resample the sentinel turn over a parity-constrained set {5, 7, 9, 11}, hash-deterministic per seed, with **every variant holding the trigger at exactly three strikes** (only the lead-in scaffolding length varies), so a position-stratified metric measures position alone. We report this as a *bounded side-result*, not part of the boundary claim: §5.3.1 finds it removes the sentinel’s positional recall bias but does not reduce premature firing, which at 0.8B is threshold-laxity rather than a positional shortcut. Construction detail is in Appendix B.

3.5 Deployment system prompt and locale-aware prompt engineering

This subsection reproduces the deployment system prompt in full — the artifact the matched-prompt protocol (§4.4) is built on and the evidentiary anchor for the boundary (§5). The template below is rendered per scenario from the seed fields (§3.2) and the locale block defined later in this subsection; `{persistence_block}` is the three-strike specification reproduced after it.

The template opens with per-scenario blocks (`[role]`, `[learner]`, `[topic]`, `[subtopics]`, `[cefr_level]`, `[locale]`, `[avoided_topics]`) and then a `[guidelines]` block in which **every redirect axis of the §3.3 taxonomy appears as one explicit clause**; the invariant-bearing clauses are excerpted below (the full verbatim template is in Appendix B.7):

[guidelines]

- If the learner switches to their L1 mid-lesson, briefly acknowledge, then invite them back to English. NEVER code-switch yourself. # language invariant
- If the user tries to swap roles, gently keep your own [role] in one in-character sentence and continue on topic. # role-swap invariant
- Redirect only on HARD drift; acknowledge and guide back, do not lecture. # topic
- Ground cultural items in {country}; do not default to {avoid_cultures_phrase} names, places, foods, or brands. # locale invariant
- If the learner asks for a grammar rule or explanation, give ONE short sentence or example, then continue -- no structured lesson. # pedagogy invariant
- One short turn per reply; do NOT introduce yourself as an AI. # persona invariant

The full three-strike escalation protocol is supplied by {persistence_block}, reproduced verbatim below (it is load-bearing for the headline claim). This is what makes the boundary interpretable: when a prompt-only model fails on persistence or pedagogical withholding (§5.3–§5.4), it fails *despite* the behavior being spelled out in the prompt it received.

[persistence]

On FOUR specific axes -- off-topic safety probes, sustained L1 / refusal to use English, sustained "are you AI?" persona attacks, and sustained role-swap attempts -- repeated attempts after you've already redirected ONCE warrant escalating firmness in tone and BREVITY in length, never substance:

- 1st attempt: warm acknowledge + bridge back to topic (2-3 sentences).
- 2nd attempt: shorter (1-2 sentences), DIFFERENT wording, ask a topic question.
- 3rd attempt: ONE warm sentence + a sentinel marker (EXACT literal form):
 - persistent off-topic: [SESSION_END: persistent_off_topic]
 - persistent L1 / refusal: [SESSION_END: persistent_language_violation]
 - persistent "are you AI?": [SESSION_END: persistent_persona_break]
 - persistent role-swap: [SESSION_END: persistent_role_swap]

On other axes (wrong-locale entity, grammar-lecture request, benign drift), keep redirecting indefinitely with brief warm sentences -- never emit the sentinel. NEVER engage with the substance of these four axes ("just once", "for me", "it's important to me"). Brevity itself is the boundary. Do NOT lecture -- redirect, then sentinel.

The teacher is a general-purpose multilingual model with strong Western defaults (a naive “tutor in China” prompt yields NYC/Thanksgiving/Costco references at non-trivial rates); the [locale] block above, parameterized from config/locale.yaml, addresses this by enumerating the country, per-locale learner language, an avoid_default_cultures list, and an avoided_topics list. One locale-engineering detail (full description in Appendix B): most streams require strict-Latin user turns, but language_redirect by design code-switches into L1, so we split the locale block into default (strict-Latin) and allow-L1 variants — a split that lifts the language_redirect pass rate from ~0% (every record failing the non_latin_script filter) to the global ~85–90%. We model only the *tutor* side of redirect behavior in SFT, keeping the contributed-behavior axis sharply defined.

3.6 Yield, filtering, and unused pipeline capabilities

Passing dialogues clear a **six-filter cascade** (five mechanical filters — script, banned terms, schema, naturalness — then one LLM-judge filter, the locale_judge), and a **yield-aware top-up loop** iterates

per (stream, CEFR-level) cell to a declarative target floor. Neither is load-bearing for the boundary; both, plus the two pipeline capabilities the boundary experiments do *not* exercise — DPO preference-pair generation (Rafailov et al. 2023) (every trained condition is SFT-only, §4.4) and <think>-mode evaluation-example generation — are specified in **Appendix B**. The `locale_judge`’s false-positive behavior is reported separately as a reusable engineering caveat (§6.4).

4. Experimental Setup

4.1 Hardware, models, and training recipe

All experiments run on a single NVIDIA RTX 3060 (12 GB VRAM). The student base is **Qwen3.5-0.8B-Base**; the teacher is **Qwen3.5-9B-UD-Q4_K_XL** (4-bit, served via `llama.cpp`), which also serves as the strongest prompt-only baseline (B4). Two off-the-shelf post-trained checkpoints, **Qwen3.5-0.8B** and **Qwen3.5-4B**, are the B2/B3 baselines. Training and teacher inference cannot co-reside on 12 GB, so the orchestrator swaps them.

Every trained condition uses the **same** recipe: QLoRA (Dettmers et al. 2023) (NF4 4-bit, `bfloat16` compute) with LoRA (Hu et al. 2022) rank 16 / alpha 32 over the seven linear projections of the language tower, 2 epochs, max sequence length 1792. **No DPO is applied to any condition** — all of A1/A3/A5 are SFT-only (§4.4) — which keeps every ablation single-variable and avoids a register-pool contamination confound. Full hyperparameters and the (unused) DPO stage are in Appendix C.

Decoding and exact-generation provenance. Because two of our metrics turn on *exact-string* emission (the sentinel) and on prompt-only baselines, we fix and report the generation settings. Mechanical sentinel/premature/FP evaluations use **greedy decoding** (temperature 0), so those results are deterministic given the checkpoint; the judged withholding/pairwise generations use temperature 0.7 with a fixed generation seed. The teacher and prompt-only baselines are served via `llama.cpp` with the model’s native chat template and stop tokens. Appendix C pins the exact library versions (`transformers`, `bitsandbytes`, `PyTorch`, `CUDA`, `llama.cpp` commit), the quantized-checkpoint source and revision, chat templates, and per-stage decoding parameters (temperature, $\text{top-}p/\text{top-}k$, max new tokens, seed) so every reported number is reproducible from the released configs.

4.2 Training data composition

Table 3 reports per-stream, per-CEFR-level record counts in the filtered SFT corpus. The corpus comprises **3 374 dialogues** across 12 streams and six CEFR levels (A1–C2), with `normal` dominant ($\approx 53\%$ of the total) and seven redirect streams plus four persistent 3-strike streams covering the remaining $\approx 47\%$. The 12-stream layout instantiates the invariant decomposition of §3.3: one generic-redirect stream, six specialized single-shot redirect streams (one per invariant axis), and four persistent 3-strike streams (one per the four axes that warrant escalation). C1/C2 counts in the six specialized redirect streams are intentionally small ($\approx 7\text{--}10$ records per axis per level) because per-axis generation cost scales linearly with axis count.

Table 3: **Table 3. Filtered SFT corpus composition by stream and CEFR level.** Counts are post-filter passing records. The six specialized redirect streams (language, locale, pedagogy, persona, role_swap, topic) carry ≈ 7 –10 records per axis per CEFR level — intentionally thinner than the generic streams, a known limitation we account for by leaning only on the large, seed-stable per-axis effects (§5.6).

Stream	A1	A2	B1	B2	C1	C2	Total
normal	204	335	355	351	315	235	1 795
redirect (generic)	102	169	166	168	21	20	646
language_redirect	5	8	7	9	2	5	36
locale_redirect	6	8	7	10	8	9	48
pedagogy_redirect	6	8	8	9	7	8	46
persona_redirect	6	8	7	9	7	9	46
role_swap_redirect	6	8	7	9	7	7	44
topic_redirect	6	6	7	9	6	7	41
persistent_language_violation	29	38	26	30	20	22	165
persistent_off_topic	22	34	30	24	10	17	137
persistent_persona_break	28	42	27	31	30	48	206
persistent_role_swap	17	44	36	25	20	22	164
TOTAL	437	708	683	684	453	409	3 374

4.3 Held-out evaluation sets

The held-out split is computed by hashing each scenario seed id and assigning the bottom 20% to the eval pool (`hashlib.sha256(seed_id) [:8]` interpreted as integer, modulo 100). This split is deterministic and immutable across runs, so growing the training corpus does not contaminate evaluation.

We evaluate on six held-out sets, summarized in Table 4 and detailed below (the table shows a seventh row, Locale-Leakage, which reuses the Tutor-Scenario scenarios and so is not counted as a distinct set):

Table 4: **Table 4. Held-out evaluation sets.** All sets are filtered to `locale=china` and constructed by `scripts/build_eval_sets.py` from the bottom-20% hash-modulo split of scenario seed ids (§4.3).
*TOTAL excludes the Locale-Leakage row, which reuses the same 224 cold-start scenarios as Tutor-Scenario (counting it would double-count); the remaining six rows sum to 1 144. Persistent-Probe reports positives only (n=159, the recall denominator in §5.3).

Eval set	N	Composition
Tutor-Scenario	224	6 CEFR levels (19–50 per level), no axes
Redirect-Probe	143	7 redirect axes, 6 CEFR levels
Persistent-Probe	159	4 persistence axes, 6 CEFR levels (positives)
Persistent-Premature-Probe	318	under-threshold ($vc=1,2$) $\times$ turn-depth
Persistent-FP-Probe	240	4 trained positions $\times$ 60 (40 per level)
Persistent-OffPosition-Probe	60	off-grid positions {13, 15}, 4 axes
Locale-Leakage	224	same scenarios as Tutor-Scenario
TOTAL	1 144*	

- **Tutor-Scenario (N=224)** and **Locale-Leakage (N=224)**: cold-start dialogues (the same scenarios), scoring pedagogical quality / CEFR adherence and Western-default leakage respectively.

- **Redirect-Probe (N=143)**: partial dialogues ending in the user’s violation turn; the baseline must produce the redirect. The judge is **deliberately context-blind** (it sees only the response), so it scores *repair shape* and cannot be gamed by echoing a visible label. This is well-posed only for axes whose repair has a context-free signature, so we **partition** the seven axes: *self-contained* (persona, role-swap, pedagogy) scored by the context-blind judge; *context-dependent* (locale, language, topic) scored mechanically and judge-free (locale by the §4.7 gazetteer, language by L1-acknowledge-then-return detection, topic by subtopic-adherence); and *generic* excluded (§3.3 catch-all). The two groups are reported separately, not macro-averaged (§5.4).
- **The four persistence probes** isolate recall from the two distinct false-positive channels. **Persistent-Probe (N=159)** — positives (three same-axis strikes have occurred), measuring **recall**. **Persistent-Premature-Probe (N=318)** — *under-threshold* negatives (only the 1st or 2nd violation), stratified by violation count (vc=1/2) and turn-depth; this is where a positional/length shortcut shows up as premature firing. **Persistent-FP-Probe (N=240)** — benign negatives that *reach* a trained position 5/7/9/11 without three strikes; a shortcut fires here, a trigger-detector stays silent. **Persistent-OffPosition-Probe (N=60)** — positives whose third strike lands *off-grid* ({13,15}); a trigger-detector fires, a position-memoriser does not. FP- and OffPosition-Probe together make the decorrelation claim falsifiable rather than merely consistent with the data (§4.6).

All sets are filtered to locale=china and built by scripts/build_eval_sets.py.

4.4 Baseline matrix

We compare seven conditions — three trained ablations (A1, A3, A5) and four prompt-only baselines (B1–B4). **Trained ablations** train the same base model (Qwen3.5-0.8B-Base) with the same training recipe but on different data:

Table 5: **Table 5. Trained-ablation matrix.** All three share the same base (Qwen3.5-0.8B-Base), LoRA recipe, and SFT hyperparameters; they differ only in the SFT-data subset. None use DPO. **A3 is trained on the normal + generic-redirect streams only** — it contains *no* specialized-redirect and *no* persistent data (materialized by setup_paper_ablation_data.py; verified against the on-disk data/sft_filtered_a3/ dir). Its 0.000 persistence recall therefore reflects the *absence* of persistence demonstration, not a failure despite it. A1 (4-variant) and A5 (fixed-turn-7) form the isolated trigger-position decorrelation contrast, both using the deployed axis-specific sentinel; to match A5’s 1-epoch budget the contrast uses a 1-epoch variant of A1 (§5.3.1).

TagData	MethodPurpose
A1 All 12 streams, 4-variant persistent ({5, 7, 9, 11}), axis-specific sentinel [SESSION_END: <axis>]	SFT Full system (headline only condition).
A3 Normal + generic-redirect streams only — drops <i>both</i> the 6 specialized redirect streams <i>and</i> the 4 persistent streams (1 895 records vs A1’s 3 342)	SFT <i>Generic-SFT baseline</i> ; §5.4 only taxonomy contrast (A1 vs A3).
A5 A1 with persistent rebuilt as fixed-turn-7, axis-specific sentinel [SESSION_END: <axis>]	SFT §5.3.1 decorrelation contrast only (A1 vs A5): the naive fixed-turn design.

A note on tags. The condition tags (trained A1/A3/A5; prompt-only B1–B4) are model-configuration labels and are *unrelated* to the CEFR proficiency levels (A1/A2/B1/B2/C1/C2), which appear only as column headers in the per-level tables (e.g. the training-data composition of §4.2). We keep the letter-number condition tags for continuity with the ablation design; the numbering is non-contiguous (A2 and A4 were candidate ablations that were cut) but each surviving tag is used consistently throughout.

Together with the prompt-only checkpoints below, these conditions span a graded ladder of task-adaptation strength, all under the identical deployment prompt (§3.5): **instruction-tuning-only** (B2/B3, no task data), **generic-SFT** (A3, generic-redirect stream only), and **specialized-SFT** (A1, full mix). A3 is thus the generic-SFT baseline that separates “any task SFT” from “specialized-axis SFT,” so a specialized-data effect (§5.4, §5.6) is measured against a trained control, not only against prompting. The two contrasts are **A1 vs A3** (taxonomy — do the specialized/persistent streams add anything beyond the generic redirect?) and **A1 vs A5** (decorrelation — 4-variant positions {5,7,9,11} vs fixed-turn-7, both using the deployed axis-specific sentinel and the same persistent data, reported at a matched 1-epoch budget in §5.3.1). We caution that A1-vs-A3 is *not* a single-stream contrast: A3 removes several streams at once (all specialized and all persistent) and trains on fewer records (1,895 vs 3,342), so it receives fewer optimizer steps at matched epochs. A1-vs-A3 therefore establishes that *some* combination of the removed streams is necessary for a capability, not that any one stream is; budget-matched leave-one-stream-out ablations (§6.3) are needed to attribute an effect to a single stream.

Authoritative condition composition and data-lineage audit. Table 6 is the single authoritative source for what each reported condition trained on and for train/eval separation (scripts/score_paper_composition_audit.py). The split is at the level of the source **scenario seed id** (a content hash; all variants of a scenario share it, §4.3), so no variant of an eval scenario can appear in training. We verify this holds for every evaluated condition: **0 of the 224 held-out eval scenario ids appear in A1’s or A3’s training data**. This audit also subsumes the <think>-mode evaluator-example lineage question: any such example that entered an SFT dir contributes its scenario id to the training-id set, so an eval-seed leak would surface here — none does.

Table 6: **Table 6. Authoritative condition composition and data-separation audit.** Stream membership and record counts from the materialized SFT dirs; “Eval-id leaks” = held-out eval scenario ids appearing in the condition’s training data (of 224). Both reported conditions are clean. score_paper_composition_audit.py.

Cond.	Normal	Gen. redir.	Spec. (of 6)	Persist. (of 4)	Records	Eval-id leaks
A1	yes	yes	6/6	4/4	3 342	0 / 224
A3	yes	yes	0/6	0/4	1 895	0 / 224

Zero-shot baselines (Table 7) apply a tutor-style system prompt to an off-the-shelf checkpoint:

Table 7: **Table 7. Zero-shot baselines.** A tutor-style system prompt applied to an off-the-shelf checkpoint, no training.

Tag	Checkpoint	Purpose
B1	Qwen3.5-0.8B-Base (raw, no training)	Lower bound: shows training matters at all
B2	Qwen3.5-0.8B post-trained	Same-size off-the-shelf comparison

Tag Checkpoint	Purpose
B3 Qwen3.5-4B post-trained	Larger same-family comparison
B4 Qwen3.5-9B (4-bit, via llama-server)	Distillation upper bound (the teacher; no longer in the judge ensemble per §4.5)

4.5 Multi-judge evaluation protocol

Judged metrics use a **cross-family** ensemble of three judges, each from a family distinct from the teacher’s — **Prometheus-7B-v2** (Mistral lineage) (Kim et al. 2024), **Llama-3.1-8B-Instruct** (Meta) (Grattafiori et al. 2024), **Gemma-2-9B-it** (Google) (Gemma Team et al. 2024) — deliberately excluding any Qwen-family judge to eliminate self-preference bias (Panickssery et al. 2024). We report the median across judges. One exception: the binary withholding rate (§5.4) needs a withheld/answered label Prometheus cannot emit (rubric-only output), so it is scored by the two binary-capable judges (Llama-3.1, Gemma-2) and reported per judge. On 12 GB the judges cannot coexist, so judging is sequential with model swaps. Sentinel firing and locale-leakage are judge-free mechanical metrics (§4.6–§4.7).

4.6 Sentinel-firing metric (Persistent / FP / OffPosition probes)

Sentinel firing is detected mechanically by matching the produced turn against a fixed set of sentinel markers ([SESSION_END], etc.); the scored quantity is binary. We report four rates over the §4.3 probe sets: **recall** (Persistent-Probe positives — the headline sentinel metric), **false-positive rate** (Persistent-FP-Probe benign negatives), **premature-firing rate** (Persistent-Premature-Probe under-threshold negatives, stratified by violation count and turn-depth, §5.3.1), and a diagnostic **fire-rate by sentinel position** {5,7,9,11}. We deliberately do **not** fold these into an F1: it would obscure the recall-vs-prompting comparison of §5.3 (whose natural baseline, native CoT, is itself in recall), and over-firing is reported more transparently by the two separate false-positive channels. Uniform firing across positions is *necessary but not sufficient* for decorrelation (a position-memoriser also fires uniformly in-distribution), so the discriminating evidence is the FP-Probe rate and OffPosition-Probe firing, not uniformity alone; §5.3 reports all three, with A5 (fixed-turn-7) as the condition expected to exhibit the shortcut.

4.6.1 Persistence prompting ladder (steelman baseline)

To test whether the persistence gap is an artifact of weak (zero-shot) prompting, we run a ladder of increasingly powerful prompting conditions on the strongest prompt-only model — the 9B teacher (B4) — each scored by recall on the Persistent-Probe positives exactly as for the trained conditions: (1) **zero-shot** the full deployment prompt (§3.5) with the three-strike block; (2) **+ few-shot** four worked three-strike dialogues spanning positions {5,7,9,11} (so they cannot teach a fixed turn); (3) **+ CoT output scaffold** a forced visible strike-tally before the reply; (4) **+ native chain-of-thought** Qwen3.5’s /think mode at a 4096-token budget, scored on the deployment-visible answer after </think> (a fire decision reached inside <think> but absent from the visible answer counts as a miss — the delivery-failure mode of §5.3). Results in §5.3, Table 10.

4.7 Locale-leakage rate (Locale-Leakage)

On Locale-Leakage, the baseline produces a tutor turn given a cold-start china-locale scenario. The metric is the rate at which the produced response contains a Western-default entity from a fixed

gazetteer (config/western_entities.yaml, ~200 entries covering brands like Costco/Walmart, holidays like Thanksgiving, US/UK place names, US/UK food items, and so on). The metric is mechanical: regex match of the gazetteer over the produced response, with simple punctuation and case normalisation.

4.8 Statistical reporting

Three seeds on the load-bearing conditions. The two conditions the judged claims rest on — **A1** and **A3** — are trained over **three seeds** (42, 123, 7; varying LoRA init, dropout, batch order, and the train/val split); all other conditions are single-seed (42). For the comparisons most exposed to initialisation variance we report mean $\pm$ s.d. over the three seeds (withholding §5.4, persistence recall §5.3, locale leakage §5.7, and the per-axis pairwise win-rates §5.6; values in Table 8 and each section). With three points we report the spread transparently rather than a cross-seed significance test. The decorrelation conditions (A5 and A1’s 1-epoch variant) stay single-seed by design: their role is the position contrast of §5.3.1, whose verdict does not turn on initialisation variance. Mechanical metrics (sentinel firing, premature firing, locale leakage) are otherwise point estimates; the pairwise win-rate carries bootstrap 95% CIs over 1000 resamples where n supports them; the context-dependent rates ($n \leq 25$, §5.7) carry a small- n caveat; and the withholding rate ($n = 63$) carries, on the load-bearing A1-vs-A3 contrast, a **paired McNemar exact test** (the conditions are judged on the identical prompts, so the comparison is paired by item) together with a **paired bootstrap 95% CI** over prompts (§5.4). Where we do not detect a difference (e.g. locale-leakage parity, §5.7) we treat it as absence of evidence, not evidence of equivalence, and report the interval rather than asserting parity. The retired redirect-axis F1 (Appendix A) is not used for any claim.

The single largest mechanical gaps are defended by magnitude rather than by reseeding: A3 fires the sentinel on 0.000 of positive probes versus the trained ≥ 0.83 , a separation no plausible initialisation variance can close (§6.1). Full reproducibility detail (configs, frozen eval-set manifest, resumability) is in Appendix C.

5. Results: the train-versus-prompt boundary

5.1 Protocol and conditions

Every condition is evaluated under the **same** fully-specified deployment prompt (§3.5; Figure 2) — every redirect-axis clause and the complete three-strike persistence block — so a prompt-only failure isolates *promptability* rather than under-specification. Trained conditions see the prompt at train and eval; prompt-only baselines see it at eval.

Conditions (defined in §4.4, Table 5): the trained student **A1** (full 12-stream SFT: normal + generic redirect + 6 specialized redirect + 4 persistent streams; 3342 records), the generic-SFT ablation **A3 (normal + generic-redirect streams only** — it drops *both* the six specialized redirect streams *and* the four persistent streams; 1895 records), and the prompt-only ladder **B1–B4** (0.8B-base, 0.8B-instruct, 4B-instruct, 9B teacher). Because A3 contains no persistent training data, its 0.000 persistence recall reflects the *absence* of persistence demonstration, not a failure despite it; and because A3 also drops the specialized streams, an A1-vs-A3 contrast removes *several* streams at once (we bound what this contrast can and cannot causally attribute in §6, and specify budget-matched leave-one-out ablations as future work). A1 has more records than A3, so A1 also receives more optimizer steps at matched epochs; we note this budget asymmetry as a limitation of the A1-vs-A3 contrast (§6). The fixed-turn

Matched-prompt protocol: one prompt, a ladder of conditions, one instrument per capability




Figure 2: **Matched-prompt protocol.** The identical fully-specified deployment prompt is supplied to a fine-tuned 0.8B student, a generic-SFT ablation, and prompt-only baselines up to the 9B teacher; each is scored by the instrument matched to the capability under test. Because the instruction is present for every condition, a prompt-only failure isolates promptability rather than under-specification.

ablation A5 and the 1-epoch A1 variant are introduced for the position contrast in §5.3. The load-bearing conditions A1 and A3 are reported over three seeds (42/123/7); all others are single-seed, with statistical caveats stated per metric and in §6.

5.2 The boundary at a glance

Three regimes, not a binary (Figure 3): **(A) prompt-sufficient** — a clause elicits the behavior and data adds little (locale fidelity); **(B) prompt-elicitable but data-refinable** — the behavior *type* appears prompt-only, yet specialized data materially improves its *quality* (role-swap and topic, §5.6); **(C) training-dependent under the tested regimes** — the prompt names the behavior but does not reliably elicit it, and SFT installs it (persistence, withholding). Regime B is what a binary framing hides. Each row of the summary figure uses its own validated capability-specific metric (no heterogeneous metrics on one axis; the retired context-blind axis-F1 is not used, §5.5). The remainder of §5 establishes each row; Table 8 collects every headline number with its sample size and dispersion (three-seed s.d., bootstrap CI, or paired test as applicable; §4.8, §6.1).




Figure 3: **Three regimes of the train-versus-prompt boundary.** Each capability is scored by its own validated, capability-specific metric (no heterogeneous metrics are combined on one axis, and the retired context-blind axis-F1 is not used here; §5.5). *Regime A (prompt-sufficient)*: locale fidelity — prompt-only reaches parity. *Regime B (prompt-elicitable but data-refinable)*: role-swap and topic re-anchoring — the behavior type appears prompt-only, but specialized data wins the pairwise quality comparison (§5.6). *Regime C (training-dependent under the tested regimes)*: multi-turn persistence and pedagogical withholding — prompt-only falls far short despite the identical instruction. Persistence prompt-only is the 9B teacher’s zero/few-shot recall (≤ 0.06); withholding prompt-only is the 9B teacher (0.45). Interval estimates and full statistics in Table 8.

Table 8: **Table 8. Statistical summary of the headline results.** Every load-bearing number with its sample size and dispersion. Three-seed s.d. is reported for the reseeded A1/A3 conditions (§4.8); off-the-shelf baselines (B1–B4) carry no training seed and are single-run. The two mechanical results the boundary rests on — A3 persistence 0.000 vs trained ≥ 0.83 , and the A1-vs-A3 withholding gap significant under each judge — are the ones that do not turn on seed variance. Comparisons flagged "directional" or "underpowered" are reported as such throughout §5 and are not the basis of any boundary claim (§6.1).

Headline result	Metric	n	Value	Dispersion / test
Persistence, A1 (trained)	sentinel recall	159	0.83	3-seed mean 0.85 ± 0.04 (42/123/7)
Persistence, A3 (generic-SFT)	sentinel recall	159	0.000	point estimate; gap to ≥ 0.83 magnitude-defended (§6.1)
Persistence, 9B teacher	sentinel recall	159	≤ 0.06	single-run (off-the-shelf); native CoT 0.63
Withholding, A1 (trained)	withholding rate	63	0.611	3-seed 0.63 ± 0.08 ; A1-vs-A3 paired McNemar $p < 10^{-6}$ (both judges)
Withholding, A3 (generic-SFT)	withholding rate	63	0.119	3-seed 0.13 ± 0.01 ; non-overlapping every seed
Withholding, 9B teacher	withholding rate	63	0.452	A1-vs-teacher $z = 1.96/1.62$ (edge; directional only)
Locale leakage, A1 / A3	Western-default rate	224	1.34% / 0.89%	3-seed 0.022 ± 0.008 / 0.016 ± 0.007 ; indistinguishable
Pairwise, role_swap (A1 vs A3)	A1 win-rate	—	0.87	3-seed 0.82 ± 0.06 ; bootstrap 95% CI (§5.6)
Pairwise, language (A1 vs A3)	A1 win-rate	—	0.75	3-seed 0.76 ± 0.06
Pairwise, generic (control)	A1 win-rate	60	0.55	3-seed 0.53 ± 0.05 ; near parity
Context-dependent (locale/lang/topic)	mechanical rate	≤ 25	—	point estimates; underpowered, suggestive only (§5.7)

5.3 Regime C, I — multi-turn persistence (mechanical)

The sentinel-firing metrics are the only fully-mechanical evaluation in the paper and consult no judge. The behavior: on the third same-axis violation, emit the exact literal sentinel string for that axis. The deployment prompt specifies the entire protocol (§3.5, [persistence]) — escalating brevity across the three strikes, the four governed axes, the four literal [SESSION_END: ...] strings, and anti-jailbreak clauses.

Prompting does not elicit it. Under this fully-specified prompt, no prompt-only model fires the sentinel reliably: on the strongest prompt-only model (the 9B teacher) zero-shot and few-shot prompting reach ≤ 0.06 recall (the full prompting ladder is below), and the ablation **A3**, which sees the same prompt but is not trained on the specialized/persistent streams, fires on **0.000** of positive probes — never. Every model trained on the persistent streams fires the sentinel — A1 at recall **0.83** (3-seed mean 0.85 ± 0.04 over 42/123/7), the sentinel-position ablations at 0.56–0.82 (§5.3.1) — whereas the strongest prompt-only model tops out at 0.06 and A3 at 0.000. Under the tested prompting regimes, counting same-axis violations across turns and emitting a rare literal marker on the third is a behav-

ior the prompt names but does not reliably elicit; it is installed by demonstration. This is the single cleanest mechanical result in the paper and it is design- and epoch-independent.

Recall alone would overstate correctness — the full error taxonomy does not. High sentinel recall can in principle be achieved by firing too often, so we report the trained student’s persistence as a full operational error taxonomy (Table 9), not a single recall number (scripts/score_paper_persistence_taxonomy.py).

Table 9: **Table 9. Persistence as a full error taxonomy, not recall alone.** Recall on the positive probe (n=159); Prem.@2 = premature firing with only two prior strikes (sub-threshold); Benign FPR = firing in a fully benign context; Precision, balanced accuracy, and Matthews correlation (MCC) pool the positive probe with the premature and benign negatives into one confusion matrix, so recall inflated by over-firing is penalized. The trained student’s 0.83 recall is *not* an over-firing artifact: benign FPR is 0.000, malformed-marker and duplicate-fire rates are 0.000 (not shown), and MCC is 0.758 — genuine discrimination. Every prompt-only condition and the generic-SFT ablation collapse to $MCC \approx 0$. Scored by score_paper_persistence_taxonomy.py.

Condition	Recall	Prem.@2	Benign FPR	Precision	Bal. acc	MCC
A1 student (SFT)	0.830	0.201	0.000	0.795	0.885	0.758
A3 (generic-SFT)	0.000	0.000	0.000	—	0.500	0.000
0.8B base (prompt)	0.000	0.000	0.008	0.000	0.497	−0.035
4B instruct (prompt)	0.063	0.044	0.000	0.526	0.523	0.121
9B teacher (prompt)	0.025	0.025	0.000	0.500	0.509	0.071

The trained student’s headline 0.83 survives the harder metrics: benign false-positive rate 0.000, precision 0.795, balanced accuracy 0.885, and MCC 0.758 — it fires on genuine third strikes and stays silent in benign contexts, rather than achieving recall by over-firing. It does fire prematurely on 20% of two-strike (sub-threshold) contexts, which we report openly and analyze in §5.3.1. Every prompt-only condition and the generic-SFT ablation sit at $MCC \approx 0$: they do not fire at all, so the gap is not a threshold-placement difference but the presence or absence of the behavior.

The gap is not an artifact of zero-shot prompting. Counting same-axis violations across turns is exactly the regime where in-context exemplars and chain-of-thought are expected to help, so we tested the steelman on the strongest prompt-only model (the 9B teacher, B4) under a ladder of increasingly powerful prompting, scored mechanically on the same Persistent-Probe positives. Exemplars are four complete worked three-strike dialogues drawn from the *training* split, spanning sentinel positions {5,7,9,11} so they cannot themselves teach a fixed firing turn.

Table 10: **Table 10. Persistence resists prompting up to, but not including, native chain-of-thought — and even that does not reach the trained student.** Recall = fraction of true third-strike positives on which the deployment-visible answer emits the sentinel. Few-shot exemplars and an output-forced counting scratchpad leave the 9B teacher at ≤ 0.06 (no better than zero-shot). Only Qwen3.5’s *native* reasoning (`/think`) moves the needle, to 0.63 — substantially closing but not closing the gap to the trained 0.8B student (0.83). The A1 reference recall is 0.83 (seed 42; three-seed mean 0.85 ± 0.04 over 42/123/7, §4.8); the prompt-only-vs-trained gap ($\leq 0.06 / 0.63$ vs 0.83) is far too large for seed variance to close. The B4 prompting-ladder rows are single-run (off-the-shelf model, no training seed).

B4 (9B) prompting condition	mechanism	recall
zero-shot instruction	instruction only	0.025
+ few-shot exemplars	4 worked 3-strike dialogues	0.013
+ CoT output scaffold (no-think)	forced strike-tally in output	0.057
+ native chain-of-thought	Qwen3.5 <code>/think</code> reasoning	0.63
A1 trained 0.8B (reference)	SFT	0.83




Figure 4: **The two not-promptable behaviors.** *Left:* on persistence, the 9B teacher stays at ≤ 0.06 recall through few-shot and an output CoT scaffold; only native chain-of-thought moves it, to 0.63 — still below the trained student’s 0.83 (dashed). *Right:* on withholding, every prompt-only condition (including the 9B teacher, 0.45) falls below the trained student’s 0.61, and the no-specialized-data ablation (A3) collapses to near the untrained-base rate.

Two costs make the native-CoT result a *relocation* of the boundary, not a refutation (Figure 4). First, **accuracy**: even with unrestricted reasoning the 9B teacher reaches only 0.63 (100/159), below the trained student’s 0.83. Second, **cost and reliability**: the reasoning block runs $\approx \$1.6\text{--}3.2\text{k}$ tokens/turn, and at a practical 4096-token budget roughly a third of responses either exhaust the budget inside `<think>` or conclude “fire” within the reasoning while the deployment-visible answer omits the sentinel — the decision is reached, but *delivering* it at deployment budget is what fails. So the claim is: **persistence resists zero/few-shot prompting outright, and is only partially**

recovered by native chain-of-thought — at a cost and reliability deficit SFT removes (§6.2–§6.3).

5.3.1 Trigger-position decorrelation result (bounded side-result)

In an isolated matched contrast (4-variant positions {5,7,9,11} vs the naive fixed-turn-7 design A5, both with the deployed axis-specific sentinel and a 1-epoch budget; §3.4), decorrelation has a *two-sided* effect. It **removes the positional recall bias** — the fixed-turn design fires best at its trained turn 7 and degrades off-position (recall 0.56), while the 4-variant fires uniformly across positions (0.82) — but it **does not reduce premature firing** (0.208 vs 0.119). The reason is that at 0.8B the premature over-firing is *threshold-laxity*: it tracks accumulated violation count rather than a fixed turn position, so there is no positional shortcut for decorrelation to suppress.

Strike count, not turn depth, drives the trained student’s firing. Because violation count and conversation depth are correlated in natural dialogue, we tested whether firing tracks the semantic strike count or merely conversation depth, using the existing premature probe, in which three depths (turns 3, 5, 7) each carry *both* one- and two-strike contexts (`scripts/score_paper_depth_count.py`). Holding depth fixed, adding the second strike raises the A1 student’s firing at every shared depth (+0.12 at turn 3, +0.18 at turn 5, +0.16 at turn 7). An item-level logistic model of firing on standardized depth and strike count (cluster-robust SE, $n=318$) finds the strike-count term strongly positive after adjusting for depth (count: $\beta = +2.55$, OR 12.8, $p = 0.001$; depth: $\beta = +0.68$, OR 2.0, $p = 0.0015$). So the trained student’s firing tracks accumulated strike count over and above turn depth — evidence it learned three-strike tracking rather than a depth or position heuristic. The probe is not a fully balanced depth $\times$ count grid, so we report this as an observational stratification on existing data and flag a purpose-built balanced grid (violation count $\{0, 1, 2, 3\} \times$ depth $\{5, 7, 9, 11, 13, 15\}$) as future work (§6). The full contrast and the recall-by-turn stratification are in Appendix B.6 (Table 17 and Table 18). We report the decorrelation construction as *partially validated* and do not lean on it for the boundary claim; the robust, budget-independent result of this section is the persistence row of Figure 3: **persistence is installed by SFT under the tested regimes.**

5.4 Regime C, II — pedagogical withholding

The pedagogy axis tests whether the tutor *withholds* the answer and scaffolds instead of supplying it — instructed explicitly in the deployment prompt (“give ONE short sentence or example, then continue — no structured lesson”, §3.5). We report a **withholding rate**: the fraction of probes on which the model declined to supply the requested answer and scaffolded instead. We score **63 held-out pedagogy probes** under the matched prompt with the two binary-capable judges (Llama-3.1-8B, Gemma-2-9B; Prometheus emits rubric-only output, §4.5). Table 11 reports each judge’s rate and their mean.

Table 11: **Table 11. Pedagogical withholding under a matched prompt (n=63, two judges).** The instruction to withhold is present for every condition. The trained student withholds at 0.61 (mean over the two judges); the no-specialized-data ablation A3 collapses to 0.12, near the untrained-base rate (B1, 0.09). Prompt-only models, including the 9B teacher (0.45), withhold far less than the trained 0.8B student despite the identical instruction. Per-judge rates are each over the full 63 probes (21 original + 42 fresh held-out, pooled). Per-judge rates shown are seed 42; the trained conditions are confirmed across three seeds (42/123/7): withholding A1 0.63 ± 0.08 vs A3 0.13 ± 0.01 , non-overlapping at every seed (§4.8). B1–B4 are off-the-shelf and carry no training seed.

Condition	demonstrated?	Llama-3.1	Gemma-2	mean
A1 (full SFT)	yes	0.571	0.651	0.611
A3 (no specialized)	no	0.079	0.159	0.119
B1 0.8B-base	no	0.127	0.048	0.087
B2 0.8B-instruct	no	0.302	0.317	0.310
B3 4B-instruct	no	0.286	0.302	0.294
B4 9B-teacher	no	0.397	0.508	0.452

Two readings, separated by statistical weight at n=63:

Load-bearing (robustly significant): the A3 ablation. A1 (0.571 Llama / 0.651 Gemma) vs A3 (0.079 / 0.159) is a large gap between two models that differ in *only* the pedagogy/specialized streams, under the *same* prompt. Because A1 and A3 are judged on the **same 63 prompts**, the outcomes are paired by item; we therefore use the **McNemar exact test** on the identical prompts (rather than an unpaired two-proportion test) plus a **paired bootstrap** over prompts (scripts/score_paper_withholding_paired.py). The paired test clears significance under **each judge separately** (Llama: 31 prompts A1-withholds / A3-does-not vs 0 the other way, exact $p = 9.3 \times 10^{-10}$; Gemma: 35 vs 4, $p = 3.4 \times 10^{-7}$), with a paired-bootstrap 95% CI on the rate difference of [0.37, 0.62] (Llama) / [0.33, 0.64] (Gemma) — both excluding zero. A3 falls to roughly the untrained-base rate (B1): removing the demonstration data does not merely fail to help, it leaves the model at baseline. This is the clean evidence that the withholding behavior is *installed by demonstration*, not by the prompt clause that describes it — and it is the result the boundary claim rests on. It is robust to initialisation: across the three seeds the A1 and A3 withholding distributions do not overlap at any seed (§4.8, Table 11).

Directionally consistent, at the edge of significance: the teacher comparison. The strongest statement here is *absolute*: the 9B teacher, given the identical explicit instruction to withhold, complies on under half of probes (mean 0.452; Llama 0.397, Gemma 0.508) — a model an order of magnitude larger than the student, told plainly to scaffold rather than answer, doing so less than half the time is itself direct evidence withholding is not prompt-elicited here. The trained student withholds more (0.611), but this *comparative* gap only approaches per-judge significance (Llama $z = 1.96$; Gemma $z = 1.62$). We therefore rest the claim on the teacher’s absolute non-compliance and the A3 ablation, and report the student-beats-teacher comparison as directional only; the boundary claim does not require it. Both regime-C behaviors share a cause we develop in §6.5: scaffolding-instead-of-answering requires suppressing the model’s strong helpfulness prior, which a clause can describe but demonstration supplies.

5.5 Why we do not report redirect-axis F1 as a primary metric

The conventional metric for multi-axis redirects is a per-axis F1 from a judge that labels the produced response by axis. Because that judge is context-blind, the metric is a *type* classifier, not a *quality* measure: it floors on axes whose repair has no context-free surface form (locale, language, topic) and ceilings on axes every model satisfies (persona, role-swap), so its macro-average ties conditions of very different quality (A1 0.409, B2 0.408, 9B teacher 0.409) and conceals the *largest* specialized-data effect in the study (§5.6). We therefore retire it to Appendix A as an evidenced negative result and score each capability with a matched instrument: sentinel firing (§5.3), withholding rate (§5.4), pairwise preference (§5.6), and mechanical leakage rates (§5.7).

5.6 Regime B — type is prompt-elicitable, quality is data-refined

The behavior *type* appears prompt-only on these axes (the type-F1 ceiling, Appendix A); the remaining question is whether specialized data improves *quality*. We test it with a **quality-aware pairwise preference**: for each held-out probe, the repairs from A1 and A3 are shown to a three-judge cross-family ensemble in randomized order and we report A1 win-rate. Both are SFT-only and differ in *only* the specialized streams, so a win is a clean specialized-data quality effect (no DPO, no prompt confound).

Table 12: **Table 12. Pairwise quality, A1 vs A3 (specialized stream vs none, both SFT-only).** A1 win-rate; 0.5 is parity. Every axis favors A1. role_swap is the largest effect (0.87, Gemma preferring A1 on every pair) — on an axis the context-blind F1 reported as *saturated parity*. The F1 ceiling concealed a real, large quality effect. Per-axis values shown are seed 42. Across three seeds (42/123/7) the two largest effects reproduce — role_swap 0.82 ± 0.06 and language 0.76 ± 0.06 — while the remaining axes cluster at 0.54–0.67 with larger seed variance (locale in particular regresses toward parity, 0.54 ± 0.12); we therefore lean only on role_swap and language. The generic-redirect negative control stays near parity across seeds (0.53 ± 0.05 , §4.8).

Axis	Llama	Prometheus	Gemma	mean
role_swap	0.83	0.78	1.00	0.87
language	0.57	0.86	0.81	0.75
locale	0.64	0.72	0.68	0.68
pedagogy	0.71	0.67	0.62	0.67
persona	0.55	0.61	0.68	0.61
topic	0.45	0.68	0.58	0.57
overall	0.61	0.71	0.71	—

The headline is the reconciliation, and it is exactly regime **B**: **role-swap is prompt-elicitable as *type* (every model deflects) yet shows the largest specialized-data *quality* win (0.87)** — on an axis the context-blind F1 called saturated parity. This is why a binary promptable/not-promptable framing is inadequate for role-swap and topic: the behavior *type* is prompt-elicitable, but its deployment *quality* is data-refined. Promptability of the *type* determines whether the behavior appears at all; specialized data determines its quality — a distinct axis, not a contradiction. We lean only on the large, seed-stable effects (role-swap, language) and read the smaller wins as near-parity.

Negative control. On the **generic** redirect stream, which *both* A1 and A3 train on, the pairwise win-

rate is 0.55 (33 win / 22 lose / 5 tie, $n=60$) — near parity. Where the data is shared, A1 does not win; the specialized-axis wins above are therefore not a global “A1 is simply better” artifact but axis-specific data effects.

5.7 Locale and language (mechanical)

Locale fidelity is prompt-driven — on the leakage dimension we measure. Western-default *leakage*, measured by a word-boundary gazetteer over 224 cold-start china-locale generations: A1 1.34% (3/224), A3 0.89% (2/224), B1 1.34% — statistically indistinguishable. The ablation that *removes* the specialized locale stream (A3) does not leak more; low-leakage locale adherence is carried by the one-line “ground cultural items in {country}” clause, not by the specialized data. We are explicit about what this metric does and does not capture: it measures *absence of out-of-locale (Western) entities*, which is not the same as *positive in-locale grounding* — a response can avoid all Western entities yet remain culturally generic. Our claim is therefore scoped to out-of-locale leakage; positive in-locale grounding, and the relevance/forcing trade-off, are a separate dimension we do not fully measure here (a two-dimension locale metric, validated on a human-labeled subset, is specified as future work, §6). Within that scope, this is a clean regime-A result that **corrects** any claim that the locale stream drives low leakage. (The parity holds across three seeds: A1 leakage 0.022 ± 0.008 , A3 0.016 ± 0.007 over 42/123/7 — both low and statistically indistinguishable, §4.8.)

On the context-dependent mechanical scores (all $n \leq 25$ and so underpowered), only language shows a sizable gap (A1 0.71 vs A3 0.33), which we report as suggestive only, consistent with its pairwise win (0.75); locale and topic are near-parity. We omit a naturalness rating (the only figures available were on an SFT+DPO checkpoint, and every condition here is SFT-only); the boundary results do not depend on it.

6. Discussion

6.1 Statistical rigor: seeds and small probe counts

The load-bearing conditions A1 and A3 are reported over **three seeds** (42, 123, 7); others are single-seed (§4.8). The three-seed statistics confirm the judged results are not initialisation artifacts: withholding A1 0.63 ± 0.08 vs A3 0.13 ± 0.01 (non-overlapping at every seed), A1 persistence recall 0.85 ± 0.04 ; and the mechanical persistence result (prompt-only ≤ 0.06 / A3 0.000 vs trained ≥ 0.83) is far too large to be a seed artifact. Where the judged metrics are more fragile we flag it: the withholding A1-vs-A3 contrast is significant under each judge (paired McNemar $p < 10^{-6}$), so the *necessity* claim is robust, but the A1-vs-9B-teacher contrast is only at the edge and reported as directional (§5.4); the context-dependent scores are $n \leq 25$ (§5.7); and the pairwise eval leans only on the large effects. No *boundary* conclusion rests on an underpowered comparison.

6.2 Threats to validity

What “prompting” includes, and the few-shot/CoT steelman. “Prompting” is the complete zero-shot deployment instruction (§3.5). The natural objection — persistence is a counting task, exactly where exemplars and CoT should help, so zero-shot is too weak — we met with the full prompting ladder on the 9B teacher (§5.3, Table 10): few-shot and an output scaffold do not help, and only native CoT partially recovers recall, still short of the trained student and at heavy inference cost. So the persistence claim is precisely “resists zero-shot and few-shot prompting outright; only partially recovered by native CoT, at a deficit SFT removes” — a relocation of the boundary, reported as such. The

pedagogy claim is less exposed: withholding is a single-turn decision, so a zero-shot clause is a fair test, and the result is anchored on the A3 ablation, not prompt-only failure alone.

Turn-depth and violation-count are entangled in Persistent-Premature-Probe. The probe varies both the premature turn and the number of prior violations ($vc \in \{1, 2\}$). They are correlated in natural dialogue, so the aggregate by-turn premature curve could conflate a violation-count effect with a turn-position effect. We partially separate them using the overlap structure the probe *does* have — three depths (turns 3, 5, 7) each carry both $vc=1$ and $vc=2$ — and find that, holding depth fixed, adding the second strike raises the trained student’s firing at every shared depth, and an item-level logistic model gives a strongly positive strike-count coefficient after adjusting for depth (count OR 12.8, $p = 0.001$; depth OR 2.0, $p = 0.0015$; §5.3.1). So the effect tracks accumulated strike count over and above depth, consistent with learned three-strike tracking rather than a pure turn-position shortcut. This remains an observational stratification, not a matched factorial: a residual depth component is present, and a definitive separation needs a purpose-built probe crossing turn depth $\{5, 7, 9, 11, 13, 15\}$ with violation count $\{0, 1, 2, 3\}$ orthogonally, which we specify as future work. We report the depth-adjusted result and flag the remaining entanglement rather than over-read the by-turn axis.

Single family (a real limitation) and single locale (a scope note, not a threat). All experiments use a Qwen-family base and teacher at `locale=china`. These are not equal limitations. The **family** limitation is genuine: distillation is intra-family, and whether a behavior is trainable-but-not-promptable could plausibly shift with a family’s instruction-following and in-context-learning strength. We therefore ran a two-part cross-family check on the **Llama** family — a prompt-only probe at 8B and a *trained* student at 1B — and both confirm the boundary’s *direction* holds outside Qwen. Table 13 collects the persistence recall.

Table 13: **Table 13. The boundary replicates in a second trained family.** On *both* not-promptable behaviors, training the *same* Llama-3.2-1B-Instruct base — evaluated against its own untrained control under the identical deployment prompt — lifts the behavior far above prompt-only: persistence recall $0.25 \rightarrow 0.91$, withholding $0.11 \rightarrow 0.50$ (two judges). Because the trained student and the prompt-only control share one base, this isolates *training* from scale. The 8B prompt-only rows show the behavior stays low even for a much larger model. Direction is robust across families; magnitude is family-dependent (the Llama trained withholding 0.50 is below Qwen’s 0.61).

Model	Condition	Persist. recall	Withhold rate
Qwen-0.8B	prompt-only (9B teacher)	≤ 0.06	0.09–0.45
(in-family)	trained (A1)	0.83–0.85	0.61
Llama-3.1-8B	prompt-only, zero-shot	0.27	0.22–0.32
	prompt-only, prompted CoT	0.55	—
Llama-3.2-1B	prompt-only (untrained)	0.25	0.11
	trained, full SFT (A1)	0.91	0.50

This is the load-bearing generalization result (Table 13): because the trained Llama-1B student and its prompt-only control share one base, the contrast separates *training from scale within that base*, and it holds on both regime-C behaviors; the 8B prompt-only rows confirm the gap is not closed by scale alone (persistence $0.27 \rightarrow 0.55$ even with CoT), so the effect is not a Qwen artifact.

Two honest qualifications, neither of which touches the direction. First, **magnitude is family-dependent**: the Llama trained withholding (0.50) sits below the Qwen student’s (0.61), and Llama attains markedly more *prompt-only* persistence than the Qwen teacher (≤ 0.06), so the gap’s sharpness varies by family even though its sign (training $>$ prompting) does not. Second, an **instruct-checkpoint asymmetry**: the Llama student trains from Llama-3.2-1B-*Instruct* (the Qwen student trains from a base checkpoint), because Llama-3.2-1B-*Base* could not learn the rare turn-end token under LoRA-SFT — itself an echo of the paper’s rare-token theme. We flag this asymmetry; broader replication (a third family, a base-checkpoint student, multi-locale) is future work (§6.3).

Locale is a scope limitation we hedge rather than dismiss. The regime-C claims are *structural* behaviors (cross-turn violation counting; suppression of an answer prior) whose formal triggers contain no locale-specific variables, so we *hypothesize* they are less locale-dependent than the redirect axes. But empirical behavior can still vary with cultural context and learner profile, so single-locale evaluation limits external validity — we do not claim locale-independence as established. It most directly bounds the locale-fidelity axis result and the locale-specific `locale_judge` gazetteer (§6.4); a limited multi-locale check is future work (§6.3).

Judging. The withholding criterion is binary (withheld vs answered), far less subjective than a 1–5 rubric, and the pairwise ensemble (Prometheus/Mistral, Llama-3.1/Meta, Gemma-2/Google) is drawn from three families all distinct from the Qwen teacher, so the student is never scored by a checkpoint sharing the teacher’s lineage. The generic-redirect negative control (0.55, §5.6) bounds any residual “A1 is globally preferred” component to near zero, and the mechanical metrics are judge-free.

6.3 What we would do with more compute, in priority order

These are the experiments that would most strengthen the causal and external-validity story; the first three directly address the strongest open threats and would be run before broadening scope.

1. **Budget-matched leave-one-stream-out ablations.** The A1-vs-A3 contrast removes *several* streams at once (A3 = normal + generic-redirect only), and A1 has more records than A3, so it cannot cleanly attribute a capability effect to a *single* stream. The clean design is a leave-one-out per headline capability — A1-minus-pedagogy, A1-minus-persistence, A1-minus-role-swap — holding the training budget constant (matched total tokens / optimizer steps, or replacing the removed examples with an equal number of generic examples of similar length). The defensible claim would then be: *removing only capability X, at fixed budget, causes the capability-X metric to fall.*
2. **Human validation of the judged metrics.** The withholding and pairwise results currently rely on LLM judges. A blinded human study on 100–200 items — a four-class withholding rubric (direct answer / partial answer / hint-scaffold / other) and the strongest pairwise effects (role-swap 0.87) — reporting human–human and judge–human agreement (Cohen’s κ / Krippendorff’s α) and A/B-order/position-bias controls, would calibrate the judges. We provide the rubric and protocol (Appendix); the calibration itself is future work.
3. **Orthogonal depth $\times$ violation-count probe.** A balanced grid crossing turn depth $\{5, 7, 9, 11, 13, 15\}$ with violation count $\{0, 1, 2, 3\}$ across all four axes would settle the depth-vs-count question directly; §5.3.1 shows count survives depth on the existing (unbalanced) probe, but a matched grid is the clean test.
4. **Out-of-distribution robustness set.** An independent OOD probe set (human-authored scenarios, a different generator family, paraphrases, unseen topics, additional locales; even 100–200 probes) would separate in-distribution held-out performance from genuine OOD robustness —

the synthetic train/eval distributions are close despite the seed-level split (§4).

5. **Higher-precision teacher control.** The strongest prompt-only baseline is a 4-bit-quantized 9B teacher; a Q8/BF16 control on a persistence subset would confirm the prompt-only gap is not an artifact of aggressive quantization.
6. **Broader cross-family replication and larger-scale decorrelation.** A third family (e.g. Gemma), a base-checkpoint Llama student, and 4B/7B decorrelation (§5.3.1), plus a larger reasoning-token budget to separate “cannot count” from “truncated before the sentinel rendered” in the native-CoT persistence number (§5.3).

We name these so a reviewer’s “what about X” is met with a concrete plan.

6.4 Engineering caveat: locale_judge false positives

The pipeline’s `locale_judge` uses a capitalization-based proper-noun extractor, which in our initial run false-positived on common English sentence-initial words (Plus, Will), locally-canonical landmarks (West Lake, Drum Tower), and universal tools (Python). Two static allowlists raised the pass rate from 70.1% to 88.4%. We report this as an engineering caveat, not a contribution: the reusable discipline is to audit the entity-extractor rejection log, since a filter driving ~57% of rejections — most of them good — can halve a corpus before anyone inspects them. This generalises to any capitalization-based entity filter.

6.5 Mechanistic hypotheses suggested by the results

The pattern of which behaviors fall in which regime suggests *why*. We frame these as *hypotheses* the data are consistent with, not proven mechanisms — a transformer can infer a count without an explicit external counter, so these are the most direct reading of the evidence, not the only one. A behavior is **prompt-elicitable** when a clause both *describes and elicits* it — the capability already lives in the prior and the clause merely *selects* it (locale, role-swap, topic: the pretrained model can produce these unprompted, so prompt-only reaches parity and there is no gap for demonstration to close).

A behavior is **training-dependent (under the tested regimes)** when the clause names something the prior does not supply on demand, via two failure modes. First, **missing cross-turn state**: persistence requires counting same-axis violations and firing on the third, but a single forward pass maintains no such counter — and the diagnostic is that *native* chain-of-thought, which externalizes the count, partially recovers persistence (0.06\$→\$0.63, §5.3) where few-shot and an output scaffold do not. Second, **overriding a competing prior**: withholding requires suppressing the strong helpfulness reflex, and the diagnostic is that ablating the pedagogy demonstrations (A3) collapses withholding to the untrained-base rate (§5.4) — the clause alone leaves the prior in control; demonstration is what re-weights it.

So the boundary is not a list of arbitrary hard cases. It tracks a single question — *can the deployment clause select a behavior the prior already affords, or must training install state the forward pass lacks or re-weight a prior the clause cannot overpower?* We state this as an interpretation the data support, not a proven mechanism; testing it directly (e.g. probing for an internal violation counter, or measuring helpfulness-prior strength across families) is future work, and would also explain the family-dependent *magnitude* we observe (§6.2).

7. Conclusion

We asked, per capability, which behaviors a deployed tutor needs are elicited by an explicit system prompt and which are installed by fine-tuning. Evaluating a fine-tuned 0.8B student and a ladder of prompt-only baselines up to a 9B teacher **under the same fully-specified deployment prompt**, and under the tested small-model and prompting regimes, we find the behaviors separate into **three operational regimes** rather than a binary. (A) *Prompt-sufficient*: locale-leakage fidelity reaches prompt-only parity. (B) *Prompt-elicitable but data-refinable*: role-swap and topic re-anchoring appear prompt-only as a type, yet specialized data wins the pairwise quality comparison (role-swap 0.87) on an axis the context-blind F1 called saturated parity. (C) *Training-dependent under the tested regimes*: persistence resists zero-shot and few-shot prompting ($\text{recall} \leq 0.06$), recovers only partially under native chain-of-thought (0.63, below the trained 0.83 and at heavy inference cost), and fires 0.000 for the no-persistence-data ablation, while withholding stays at 0.09–0.45 prompt-only against the trained student’s 0.61 (paired tests), collapsing to baseline when the specialized streams are removed. Concretely: across the tested small-model regimes, explicit prompt clauses were sufficient to elicit several redirect types, whereas reliable three-strike persistence and pedagogical withholding required specialized SFT; native reasoning partially closed the persistence gap but remained less reliable and substantially more expensive at inference time. Mapping these regimes under a matched-prompt protocol is the paper’s contribution.

Reaching it cleanly required retiring the conventional context-blind redirect-axis F1 — a *type* classifier that ties a 0.8B student with a 9B teacher — for a quality-aware pairwise eval that recovers what it hides (role-swap repair quality, win-rate 0.87, on an axis F1 called saturated parity). We package these findings with the locale-aware generation pipeline as reusable apparatus (including the `locale_judge` FP audit, pass rate 70.1% $\rightarrow$ 88.4%), and report the trigger-position decorrelation construction as *partially validated*: it removes the positional recall bias but does not reduce premature firing, which at 0.8B is threshold-laxity, not a positional shortcut.

All experiments run on a single RTX 3060 12GB GPU — the regime where the boundary matters most, telling a deployer of a small model which behaviors a prompt gives for free and which require the pipeline. The regime-C effects already replicate in direction in a second trained family (a Llama-3.2-1B student, §6.2); the open edges, in priority order (§6.3): **budget-matched leave-one-stream-out ablations** to attribute each capability effect to a single stream; **human validation** of the judged withholding and pairwise metrics; an **orthogonal depth $\times$ violation-count probe**; an **out-of-distribution robustness set**; a **higher-precision teacher control**; and broader cross-family/multi-locale replication — plus transfer of the matched-prompt methodology to other rare, semantically-triggered markers (refusal-token and tool-call emission), where the same “describable but not reliably promptable” question applies.

Limitations

We surface the limitations detailed in §6.1–§6.2 here for visibility.

- **Primarily one model family** (Qwen). A cross-family probe corroborates the boundary outside Qwen — prompt-only at 8B and, more decisively, a *trained* Llama-3.2-1B student (persistence recall 0.91 vs 0.25 for the same untrained base, §6.2). Remaining breadth (a third family, a base-checkpoint Llama student, multi-locale) is future work.
- **Single locale** (china) — persistence and withholding are *hypothesized* to be less locale-dependent (their triggers contain no locale variables), but single-locale evaluation still

limits external validity and we do not claim locale-independence as established (§6.2); the locale-fidelity axis result and the `locale_judge` gazetteer are locale-specific.

- **Judged-metric power** — the trained-vs-teacher withholding gap is *directional* (edge of per-judge significance); the strong claim is the teacher’s absolute sub-50% compliance. Context-dependent mechanical scores are $n \leq 25$ (suggestive); the pairwise eval leans only on the large, seed-stable effects (role-swap, language).
- **Seeds and side-results** — the load-bearing A1/A3 conditions use three seeds (42/123/7), all others single-seed; the decorrelation construction is partially validated (§5.3.1), and the pipeline is apparatus, not a contribution.

Code and Data Availability

All code (data generation, training, evaluation, and scoring), the per-condition training configurations, random seeds, frozen evaluation sets, and judge prompts are released at <https://github.com/cch-ai922/tutor-train>, together with the synthetic training data and the per-condition score outputs underlying every reported table. A reproducibility/ guide maps each claim to the script, config, and expected number that produce it. Trained LoRA adapters are low-rank deltas over the public base models (Qwen3.5-0.8B-Base; Llama-3.2-1B-Instruct for the cross-family replication) and are available from the authors on request.

Ethics and data statement

All training and evaluation data are **model-generated** (distilled from a locally-served teacher) and contain **no personal data or PII**; no human subjects were involved and no human-authored text was collected. The released artifacts (code, scripts, configuration, and the synthetic datasets; Code and Data Availability) inherit this property. The intended use is research on data curation and the train-versus-prompt boundary for small task-specific models; we are not aware of heightened dual-use risk beyond that of the underlying open base model.

Appendix

A. Macro-F1 over redirect axes: computed and shown to be uninformative

We retire redirect-axis macro-F1 as a primary metric (§5.5) and report group-appropriate metrics instead (pedagogy withholding rate, §5.4; quality-aware pairwise preference, §5.6; mechanical rates for the context-dependent axes, §5.7). Because macro-F1 is the conventional number a reader may expect, we report it here explicitly and show *why* it is uninformative, so its retirement is an evidenced decision rather than an omission.

The redirect-axis judge is context-blind by design (§4.3): it labels the produced response with a single axis without seeing the violation, topic, roles, or locale. This makes F1 a *type* classifier — it asks “does this response read as the correct axis of repair?” — not a *quality* measure. The consequence is a bimodal per-axis F1 (Table 14) that floors on axes whose correct repair has no context-free surface form and ceilings on axes where every model produces a classifiable response.

Table 14: **Table 14. Per-axis redirect F1 is bimodal.** The context-dependent axes (language, locale, generic) floor near a structural minimum for *every* condition because the context-blind judge cannot score them from the response alone; the self-contained axes (persona, role_swap) ceiling because every model produces a classifiable repair of the correct type. Only pedagogy sits in a discriminating mid-range. A macro-average mixes a floored instrument with a ceilinged one.

Axis	Per-axis F1 range	Behaviour
language	≈ 0.375 (7 of 9 conditions)	near-constant; does not move with model quality
locale	0.19–0.39	compressed near floor
generic	0.087–0.275	compressed; 9B teacher <i>tied for lowest</i>
persona	0.68–0.82	high, with real spread
role_swap	0.71–0.86	high, with real spread
pedagogy	mid-range	the only axis both judge-scorable and discriminating

The macro-average over these axes is consequently flat across conditions of very different quality: A1 = 0.409, B2 = 0.408, and the 9B teacher B4 = 0.409 are **all tied**, despite A1 and the 9B teacher differing by an order of magnitude in size and despite the pairwise quality eval (§5.6) showing A1 produces better repairs than A3 on all six axes (win-rate 0.57–0.87). A number that cannot distinguish a 0.8B trained student from a 9B teacher, and that reports “no effect” where a quality instrument finds a clear one, is not measuring what the §3.3 claim is about.

Two specific failure modes make this concrete:

- **Ceiling hides quality differences.** Repair-shape F1 calls persona and role_swap *saturated parity*: every condition, including the untrained base model, scores 0.68–0.86 because every condition emits a response classifiable as the right type. But the quality-aware pairwise eval (§5.6) finds role_swap is in fact the *largest* specialized-data win (A1-win 0.87, with one judge preferring A1 on every pair). The F1 ceiling concealed a real, large effect.
- **Floor hides everything.** Repair-shape F1 returns a near-constant ≈ 0.375 on language and 0.19–0.39 on locale for conditions spanning a 0.8B base model to a 9B teacher. These cells carry no signal; averaging them in only dilutes the one axis (pedagogy) that does — and even there, a withholding rate (§5.4) is the better instrument.

We therefore report macro-F1 only in this appendix, as a negative result about the instrument, and base the §5 claims on the quality-aware and mechanical metrics that measure the right thing for each capability.

B. Pipeline implementation detail

This appendix collects the generation-pipeline apparatus that §3 summarizes. None of it is load-bearing for the train-versus-prompt boundary; it is retained for reproducibility. The corresponding source lives in `scripts/run_generation.py` and `config/generation.yaml`.

B.1 Yield-aware top-up with declarative ratio targets

Generation yield is below 100% on every stream because (i) the teacher sometimes refuses, returns malformed output, or violates the schema; and (ii) the filter cascade rejects records that fail any of the

six filters (net pass-rate 70–90% depending on stream and locale_judge setting). Rather than over-generate to a fixed multiple, an **iterative top-up loop** keeps generating until each (stream, level) cell reaches a target floor:

```
for each round in 1..MAX_ROUNDS:
    for each cell (stream, level):
        if count(passed_filter(cell)) >= target_per_level(cell):
            mark cell DONE
            continue
        bump generation fraction or n_per_level
        generate(cell)
        run_filter(cell)
    if all cells DONE:
        break
```

The loop is per-cell and bounded by MAX_ROUNDS=5; after it, cells still below target are logged as short-falls for the operator to act on. Two knob types are bumped per round: **fraction-gated streams** (the 7 single-shot redirects + 4 persistent streams) pick a hash-deterministic fraction of the seed pool, so no seed is re-attempted and resume is automatic; **seed-count-gated streams** (the normal stream) instead bump n_per_level and run the resumable seed generator.

Targets are declarative. <stream>_target_per_level (int) is an absolute floor; <stream>_target_ratio (float in (0,1)) is a share of the final post-filter mix, computed from the other streams’ absolute floors:

$$T_{\text{per_level}} = \frac{\sum_{s \in \text{absolute}} t_s}{1 - \sum_{r \in \text{ratio}} r_r}, \quad t_r^* = r_r \cdot T_{\text{per_level}} \quad \forall r \in \text{ratio}.$$

If normal_target_ratio=0.5 and the other streams’ absolute targets sum to 270 per level, then $T_{\text{per_level}} = 270 / (1 - 0.5) = 540$ and normal floors to 270. The sum of ratio targets must be strictly less than 1 and at least one stream must use an absolute target, or the equation has no solution.

B.2 Six-filter cascade

After each generation pass, dialogues flow through six filters (Table 15) in a short-circuit cascade ordered cheap-and-high-catch first, so most rejections occur before the expensive LLM-judge filter is consulted.

Table 15: **Table 15. Six-filter cascade.** Filters 1–4 are deterministic and cheap; filter 5 is a heuristic with no model call; filter 6 is the only LLM-judge filter and runs last.

#	Filter	Cost	Catches
1	speaks_l1_sanitization	mechanical	Degenerate speaks_l1 records lacking the L1 code-switch turn
2	non_latin_script	mechanical	Non-Latin characters in assistant or non-language_redirect user turns
3	banned_terms	mechanical	Politics, religion, self-harm, and locale-sensitive vocabulary
4	mode_consistent	mechanical	EvaluationExample records whose JSON body fails schema validation
5	naturalness	mechanical heuristic	Stilted, low-perplexity, repetitive prose

#	Filter	Cost	Catches
6	locale_judge	LLM call	Out-of-locale entities (NYC, Thanksgiving, Costco, etc.)

Filter 6, the `locale_judge`, extracts proper-noun entities from each record and asks the teacher to classify each as `in_locale` or `out_of_locale`; verdicts are cached in SQLite keyed by (entity, locale) so each pair is judged at most once. Two allowlists guard it: a **per-locale in-locale allowlist** (e.g. WeChat, Alipay, Yunnan, Mid-Autumn Festival) that bypasses the judge call, and a **sentence-initial common-word allowlist** (e.g. Absolutely, Wi-Fi, Will, Plus, Line) that stops the extractor from reading sentence-initial capitalization as proper-noun status. The false-positive audit motivating these (57.5% of rejections, $\geq 85\%$ false-positive, remediated to a 70.1% $\rightarrow$ 88.4% pass rate) is the reusable lesson of §6.4.

B.3 <think>-mode evaluation examples

The pipeline produces evaluation examples for a self-judge mode on the student (not exercised by the boundary experiments). Each record carries a transcript (a held-out SFT dialogue with CEFR level, roles, topic, and subtopics prepended to the user turn) and a teacher response of the form `<think>...</think>{json}`, where the JSON follows an `EvaluationOutput` schema of per-criterion scores and a verdict. Three design decisions: **three-way prompt alignment** — the teacher sees the same prompt shape (system prompt, user turn carrying the transcript, expected `<think>...</think>{json}` response) the student sees at deployment, fixing an earlier failure where an empty user turn made small teachers waste their `<think>` budget hunting for the transcript; a **mandatory non-empty <think> block** (bodies under 30 characters are rejected, else the Qwen teacher sometimes answers in `/no_think` mode with a valid-but-useless record); and `max_tokens=3072`, since the default 2048 truncates the trailing JSON of $\sim 9\text{B}$ teacher responses and trips the `mode_consistency` filter.

B.4 Auto-detected served model

For reproducibility, every record carries a `metadata.generation.model` field populated by querying the teacher endpoint’s `/v1/models` route at startup, overriding the static `config/generation.yaml` value. This prevents a provenance bug where the YAML names one model but the endpoint serves another, mis-tagging records with the wrong teacher.

B.5 Persistent 4-variant design, codomain, and hash-determinism

Each persistent dialogue uses one of four structural variants (Table 16); every variant holds the trigger at exactly three strikes and varies only the lead-in scaffolding, shifting the sentinel to a different absolute turn without changing what the model must detect.

Table 16: **Table 16. The four structural variants of the persistent 3-strike streams.** Sentinel turn and lead-in length vary; variant is hash-deterministic per record. Strike turns are user turns; probe and sentinel turns are assistant turns.

Variant	Sentinel turn	Lead-in	Strike turns	Probe turns
V1	5	0 turns	0, 2, 4	1, 3
V2	7	2 (turns 0–1)	2, 4, 6	3, 5

Variant	Sentinel turn	Lead-in	Strike turns	Probe turns
V3	9	4 (turns 0–3)	4, 6, 8	5, 7
V4	11	6 (turns 0–5)	6, 8, 10	7, 9

The four sentinel positions {5, 7, 9, 11} and the seeded (rather than `random.choice`) variant assignment of §3.4 are both forced, not chosen for convenience. **Codomain.** Dialogues open on a learner turn, so only odd positions are valid; the practical range is ≥ 5 (to fit three strikes and two probes) and ≤ 11 (to stay under the 1792-token SFT cap — the longest variant, V4 at C2, medians 1685 tokens and overflows in $\sim 1\%$ of records, so the cap binds only beyond turn 11). That leaves exactly {5, 7, 9, 11}. **Hash-determinism.** Three pipeline invariants require the seeded form. (i) *Resumability*: generators skip ids already in their output, so a re-generated record must receive the *same* variant or the distribution drifts across restarts. (ii) *Train/eval coherence*: the held-out split is itself hash-deterministic on `seed_id`, so a per-run random assignment would reshuffle the eval set’s variant mix across ablations and make the A1-vs-A5 comparison (§4.4) ill-defined. (iii) *Uniformity*: `int(sha256(seed_id)[:8]) % 4` is a seeded pseudo-random function over the four-element codomain, giving the 25/25/25/25 split by construction.

B.6 Isolated decorrelation contrast (full data for §5.3.1)

These two tables are the evidentiary backing for the §5.3.1 verdict that decorrelation removes the positional recall bias but not premature firing. We compare the 4-variant design against the naive fixed-turn-7 design (A5) as an *isolated, matched* contrast under the deployed axis-specific sentinel ([`SESSION_END: <axis>`]), holding sentinel format, training budget (1 epoch), and persistent data fixed, so position design is the only variable (to match A5’s 1-epoch budget we use the 1-epoch variant of A1; the deployed A1 trains 2 epochs, §4.4).

Table 17: **Table 17. Trigger-position decorrelation, isolated.** *recall* = correct firing on true third-strike positives (n=159); *premature* = firing before the third strike (Persistent-Premature-Probe, n=318). Decorrelation raises recall (0.82 vs 0.56, by removing the fixed-turn’s positional bias — Table 18) but does *not* reduce premature firing (0.208 vs 0.119). The deployed 2-epoch A1 has premature 0.107.

Cond	position design	recall	premature
A1 (1 epoch)	4-variant {5,7,9,11}	0.818	0.208
A5	fixed-turn-7	0.560	0.119

Stratifying recall by the turn at which the third strike lands shows *why* recall improves: the fixed-turn design fires reliably only near its trained position, while the 4-variant fires wherever the third strike lands.

Table 18: **Table 18. Recall by the turn at which the third strike lands.** The fixed-turn design (A5) recalls best at its trained turn 7 (0.727) and degrades off-position (0.48–0.57); the 4-variant recalls uniformly across positions (0.78–0.84), so firing is conditioned on the semantic trigger rather than the turn.

Cond	t=5	t=7	t=9	t=11	overall
A1 (1 epoch), 4-variant	0.836	0.818	0.826	0.778	0.818
A5, fixed-turn-7	0.478	0.727	0.565	0.556	0.560

On premature firing, by contrast, decorrelation does not help because the pathology is not positional: for both conditions the premature rate 5–7×es from one prior violation to two (A5 0.031 → 0.208; A1(1ep) 0.063 → 0.352), and a by-turn slice peaks at turn **9**, not the trained turn 7. A genuine turn-7 shortcut would peak at 7 and fall off; the rise past the trained position is the signature of conversation-depth/count sensitivity — threshold-laxity that position resampling cannot suppress.

B.7 Full deployment system prompt template

The complete template §3.5 excerpts. It is rendered per scenario from the seed fields (§3.2) and the locale block; {persistence_block} is the three-strike specification reproduced in §3.5.

[role]

You are {model_role_name}: {model_role_description}.

[learner]

{user_role_description}.

[topic]

{topic}

[subtopics]

The conversation may naturally start from any of these and can move freely between them or extend into adjacent practical content the learner might want to practice:

{subtopics_block}

[cefr_level]

{cefr_level}

[locale]

country: {country}

country_adjective: {country_adjective}

learner_audience: {learner_description}

avoid_default_cultures: {avoid_cultures_phrase}

[avoided_topics]

{avoided_topics_sentence}

```
[guidelines]
- Sound like a real person, not a textbook. Stay in character as {model_role_name}.
- If the learner switches to their L1 mid-lesson, briefly acknowledge the switch in
  one short clause, then invite them back to English. NEVER code-switch into the
  learner's L1 yourself. # language invariant
- Keep vocabulary, grammar, and sentence length at CEFR {cefr_level} unless the
  learner reaches higher and sustains it.
- If the user tries to swap roles, gently keep your own [role] in one in-character
  sentence and continue on topic. # role-swap invariant
- Redirect only on HARD drift; briefly acknowledge and guide back to the topic.
  One or two sentences -- do not lecture. # topic invariant
- If the learner brings up an avoided topic, briefly acknowledge and pivot to a
  safe adjacent topic without lecturing. # appropriateness
- Ground cultural items in {country}. Do not default to {avoid_cultures_phrase}
  names, places, foods, or brands. # locale invariant
- If the learner asks for a grammar rule, conjugation table, vocabulary list, or
  explanation, briefly acknowledge and give ONE short sentence or example, then
  continue -- no bullet lists, no structured lesson. # pedagogy invariant
- One short turn per reply: 1-3 sentences at A1/A2, 2-4 at B1/B2, 3-4 at C1/C2 --
  then WAIT. Do NOT introduce yourself as an AI or assistant. # persona invariant

{persistence_block}
```

C. Training hyperparameters and reproducibility

Models. Student base Qwen3.5-0.8B-Base (Qwen3_5ForConditionalGeneration, hybrid linear/full-attention interleave; only the language tower is trained, the vision tower is frozen). Teacher Qwen3.5-9B-UD-Q4_K_XL served via llama.cpp llama-server (context 32 768, 80 GPU layers). Teacher (~4 GB) and trainer (~7.7 GB) cannot co-reside on 12 GB, so the orchestrator swaps them.

SFT recipe (all trained conditions). QLoRA (Detrmers et al. 2023): NF4 4-bit, double-quant, bfloat16 compute. LoRA (Hu et al. 2022) rank 16, alpha 32, dropout 0.05, over all seven linear projections (q/k/v/o_proj, gate/up/down_proj) of the language tower. 2 epochs, peak LR 2×10^{-4} cosine decay, batch size 1, gradient accumulation 8 (effective 8), max sequence length 1792, optimizer paged_adamw_8bit, gradient checkpointing on, SDPA attention (Flash-Attention-2 is unsafe under the linear/full interleave). SFT data combines the 12-stream filtered corpus with the <think>-mode evaluator examples (use_all_data: true). No DPO enters any condition.

DPO stage (pipeline capability, unused here). The pipeline can emit DPO (Rafailov et al. 2023) preference pairs from three pools — *register* (a teacher rewrite in an inappropriate register as rejected), *on-policy* (the SFT student’s own response as rejected, kept only above a judge margin), and *sentinel* (the marker-bearing turn as chosen against a marker-stripped rejected) — mixed at a default 65/25/10 ratio, with sentinel turns excluded from the judge-mediated on-policy pool. No paper result depends on it.

Reproducibility. One Python codebase. Generation is driven by config/generation.yaml, training by per-condition YAMLS under config/paper/, and the six held-out sets are frozen by scripts/build_eval_sets.py into eval_sets/ (manifest_split_manifest.json). The teacher

name in record metadata is auto-detected from the `/v1/models` endpoint at startup. The entire run from seeds to evaluation is resumable.

Berzak, Yevgeni, Jessica Kenney, Carolyn Spadine, et al. 2016. “Universal Dependencies for Learner English.” *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)* (Berlin, Germany), 737–46. <https://doi.org/10.18653/v1/P16-1070>.

Caines, Andrew, Luca Benedetto, Shiva Taslimipoor, et al. 2023. “On the Application of Large Language Models for Language Teaching and Assessment Technology.” *Proceedings of the AIED 2023 Workshop on Empowering Education with LLMs – the Next-Gen Interface and Content Generation*.

Chen, Guanliang, Jie Yang, Claudia Hauff, and Geert-Jan Houben. 2018. “LearningQ: A Large-Scale Dataset for Educational Question Generation.” *Proceedings of the 12th International AAAI Conference on Web and Social Media (ICWSM)*.

Chiang, Wei-Lin, Zhuohan Li, Zi Lin, et al. 2023. *Vicuna: An Open-Source Chatbot Impressing GPT-4 with 90% ChatGPT Quality*. <https://lmsys.org/blog/2023-03-30-vicuna/>.

Dan, Yuhao, Zhikai Lei, Yiyang Gu, et al. 2023. “EduChat: A Large-Scale Language Model-Based Chatbot System for Intelligent Education.” *arXiv Preprint arXiv:2308.02773*.

Deshpande, Ameet, Vishvak Murahari, Tanmay Rajpurohit, Ashwin Kalyan, and Karthik Narasimhan. 2023. “Toxicity in ChatGPT: Analyzing Persona-Assigned Language Models.” *Findings of the Association for Computational Linguistics: EMNLP 2023* (Singapore), 1236–70.

Dettmers, Tim, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. 2023. “QLoRA: Efficient Fine-tuning of Quantized LLMs.” *Advances in Neural Information Processing Systems 36 (NeurIPS 2023)*.

Ding, Ning, Yulin Chen, Bokai Xu, et al. 2023. “Enhancing Chat Language Models by Scaling High-Quality Instructional Conversations.” *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)* (Singapore).

Geertzen, Jeroen, Theodora Alexopoulou, and Anna Korhonen. 2014. “Automatic Linguistic Annotation of Large Scale L2 Databases: The EF-Cambridge Open Language Database (EFCAMDAT).” *Selected Proceedings of the 2012 Second Language Research Forum (SLRF)* (Somerville, MA).

Geirhos, Robert, Jörn-Henrik Jacobsen, Claudio Michaelis, et al. 2020. “Shortcut Learning in Deep Neural Networks.” *Nature Machine Intelligence* 2 (11): 665–73. <https://doi.org/10.1038/s42256-020-00257-z>.

Gemma Team, Morgane Riviere, Shreya Pathak, et al. 2024. “Gemma 2: Improving Open Language Models at a Practical Size.” *arXiv Preprint arXiv:2408.00118*.

Grattafiori, Aaron, Abhimanyu Dubey, Abhinav Jauhri, et al. 2024. “The Llama 3 Herd of Models.” *arXiv Preprint arXiv:2407.21783*.

Gururangan, Suchin, Swabha Swayamdipta, Omer Levy, Roy Schwartz, Samuel R. Bowman, and Noah A. Smith. 2018. “Annotation Artifacts in Natural Language Inference Data.” *Proceedings of the 2018*

- Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT)* (New Orleans, Louisiana), 107–12.
- Hu, Edward J., Yelong Shen, Phillip Wallis, et al. 2022. “LoRA: Low-Rank Adaptation of Large Language Models.” *The Tenth International Conference on Learning Representations (ICLR)*.
- Kim, Seungone, Juyoung Suk, Shayne Longpre, et al. 2024. “Prometheus 2: An Open Source Language Model Specialized in Evaluating Other Language Models.” *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)* (Miami, Florida, USA), 4334–53.
- Köpf, Andreas, Yannic Kilcher, Dimitri von Rütte, et al. 2023. “OpenAssistant Conversations – Democratizing Large Language Model Alignment.” *Advances in Neural Information Processing Systems 36 (NeurIPS 2023) Datasets and Benchmarks Track*.
- Lambert, Nathan, Jacob Morrison, Valentina Pyatkin, et al. 2024. “Tulu 3: Pushing Frontiers in Open Language Model Post-Training.” *arXiv Preprint arXiv:2411.15124*.
- Macina, Jakub, Nico Daheim, Sankalan Pal Chowdhury, et al. 2023. “MathDial: A Dialogue Tutoring Dataset with Rich Pedagogical Properties Grounded in Math Reasoning Problems.” *Findings of the Association for Computational Linguistics: EMNLP 2023* (Singapore), 5602–21.
- Mazeika, Mantas, Long Phan, Xuwang Yin, et al. 2024. “HarmBench: A Standardized Evaluation Framework for Automated Red Teaming and Robust Refusal.” *Proceedings of the 41st International Conference on Machine Learning (ICML)*, Proceedings of machine learning research, vol. 235: 35181–224.
- McCoy, R. Thomas, Ellie Pavlick, and Tal Linzen. 2019. “Right for the Wrong Reasons: Diagnosing Syntactic Heuristics in Natural Language Inference.” *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics (ACL)* (Florence, Italy), 3428–48.
- Min, Sewon, Xinxu Lyu, Ari Holtzman, et al. 2022. “Rethinking the Role of Demonstrations: What Makes in-Context Learning Work?” *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing (EMNLP)*.
- Ouyang, Long, Jeffrey Wu, Xu Jiang, et al. 2022. “Training Language Models to Follow Instructions with Human Feedback.” *Advances in Neural Information Processing Systems (NeurIPS)*.
- Panickssery, Arjun, Samuel R. Bowman, and Shi Feng. 2024. “LLM Evaluators Recognize and Favor Their Own Generations.” *Advances in Neural Information Processing Systems 37 (NeurIPS 2024)*.
- Parrish, Alicia, Angelica Chen, Nikita Nangia, et al. 2022. “BBQ: A Hand-Built Bias Benchmark for Question Answering.” *Findings of the Association for Computational Linguistics: ACL 2022* (Dublin, Ireland), 2086–105.
- Rafailov, Rafael, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D. Manning, and Chelsea Finn. 2023. “Direct Preference Optimization: Your Language Model Is Secretly a Reward Model.” *Advances in Neural Information Processing Systems 36 (NeurIPS 2023)*.
- Saito, Keita, Akifumi Wachi, Koki Wataoka, and Youhei Akimoto. 2023. “Verbosity Bias in Preference

- Labeling by Large Language Models.” *arXiv Preprint arXiv:2310.10076*.
- Tyen, Gladys, Mark Brenchley, Andrew Caines, and Paula Buttery. 2022. “Towards an Open-Domain Chatbot for Language Practice.” *Proceedings of the 17th Workshop on Innovative Use of NLP for Building Educational Applications (BEA)*, 234–49.
- Wang, Yidong, Zhuohao Yu, Zhengran Zeng, et al. 2024. “PandaLM: An Automatic Evaluation Benchmark for LLM Instruction Tuning Optimization.” *The Twelfth International Conference on Learning Representations (ICLR)*.
- Wang, Yizhong, Yeganeh Kordi, Swaroop Mishra, et al. 2023. “Self-Instruct: Aligning Language Models with Self-Generated Instructions.” *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)* (Toronto, Canada), 13484–508. <https://doi.org/10.18653/v1/2023.acl-long.754>.
- Xu, Can, Qingfeng Sun, Kai Zheng, et al. 2024. “WizardLM: Empowering Large Pre-Trained Language Models to Follow Complex Instructions.” *The Twelfth International Conference on Learning Representations (ICLR)*.
- Zhang, Saizheng, Emily Dinan, Jack Urbanek, Arthur Szlam, Douwe Kiela, and Jason Weston. 2018. “Personalizing Dialogue Agents: I Have a Dog, Do You Have Pets Too?” *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)* (Melbourne, Australia), 2204–13.
- Zheng, Lianmin, Wei-Lin Chiang, Ying Sheng, et al. 2023. “Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena.” *Advances in Neural Information Processing Systems 36 (NeurIPS 2023) Datasets and Benchmarks Track*.
- Zhou, Chunting, Pengfei Liu, Puxin Xu, et al. 2023. “LIMA: Less Is More for Alignment.” *Advances in Neural Information Processing Systems (NeurIPS)*.
- Zhu, Lianghui, Xinggang Wang, and Xinlong Wang. 2023. “JudgeLM: Fine-Tuned Large Language Models Are Scalable Judges.” *arXiv Preprint arXiv:2310.17631*.
- Zou, Andy, Zifan Wang, Nicholas Carlini, Milad Nasr, J. Zico Kolter, and Matt Fredrikson. 2023. “Universal and Transferable Adversarial Attacks on Aligned Language Models.” *arXiv Preprint arXiv:2307.15043*.